

# Kafka monitoring using Prometheus and Grafana

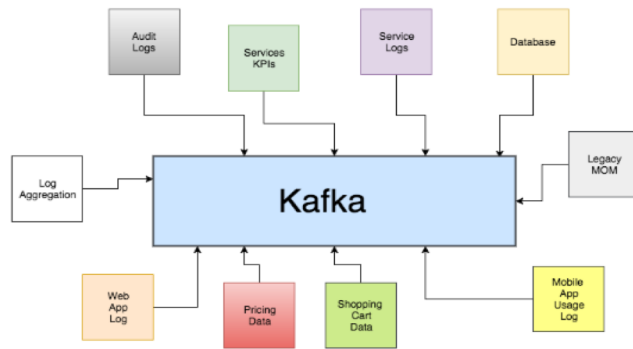
Dr. Soheila Dehghanzadeh  
Denso Automotive Deutschland

# Agenda

- Introduction to Kafka, Prometheus and Grafana
- Install Kafka
- Install Prometheus
- Monitor Kafka from Prometheus
- Install Grafana
- Monitor Kafka from Grafana
- Prometheus Rules

# Introduction to Kafka

- Distributed streaming platform
  - Publish and subscribe to streams of records
  - Fault tolerant storage
    - Replicate topic log partitions to multiple servers
  - Application can process records as they appear in kafka
  - Fast and efficient IO by batching and compressing records
- Decouples data streams (consumers do not need to know about producers)
- Ingest data into data lakes, applications and real-time analytics systems
- Usecases
  - IoT applications
  - Microservices
  - Distributed applications




# Introduction to Zookeeper

- A distributed, open-source coordination system for distributed applications
  - **Configuration** : sharing config info across all nodes
  - **Synchronization: locks, barriers and queus**
  - **Naming registry** find a machine in a cluster of 1000 of services
  - **Group services**: leader election

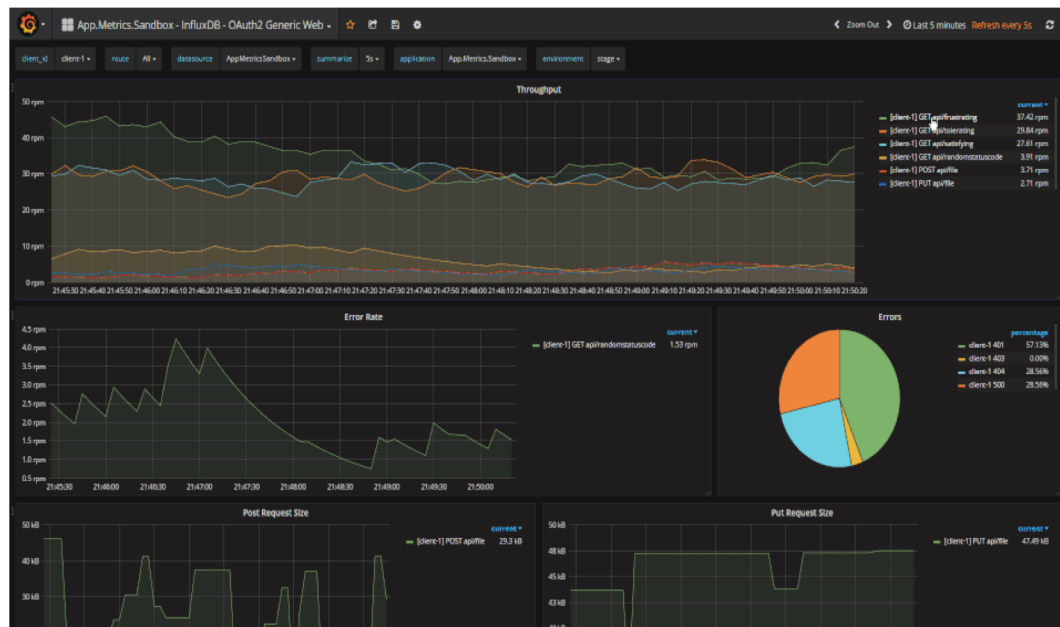
# Introduction to Prometheus

- Prometheus is a metrics-based time series database
- Metric types
  - Counter: an increasing value
  - Gauge: an arbitrary value that can go up or down
  - Histogram: sampled observations counted in buckets
- Prometheus consists of two parts:
  - An exporter exposing metrics at a port
    - Node exporter, MySQL exporter,...
  - A Prometheus server



# Introduction to Grafana

- Receive time-series data and visualize it with customized dashboards
  - Imports from
    - [Graphite](#)
    - [Prometheus](#)
    - [InfluxDB](#)
    - [Elasticsearch](#)
    - [Google Stackdriver](#)
    - [AWS CloudWatch](#)
    - [Azure Monitor](#)
    - [Loki](#)
    - [MySQL](#)
    - [PostgreSQL](#)
    - [Microsoft SQL Server \(MSSQL\)](#)
    - [OpenTSDB](#)
    - [Testdata](#)
    - [Mixed](#)



# Scenario

1. In this scenario, we aim to run a kafka cluster and monitor it via
  - a. Prometheus
  - b. Grafana
2. To continue with the handson create a folder to keep the tutorial materials
  - a. `mkdir summerSchool`
3. Prerequisites
  - a. Linux (Debian/Ubuntu/Mint)
  - b. Java

`sudo apt update`

`sudo apt install default-jdk`

# Install Kafka

## 1. Install kafka from <https://tecadmin.net/install-apache-kafka-ubuntu/>

- a. `wget http://www-us.apache.org/dist/kafka/2.2.1/kafka_2.12-2.2.1.tgz`
- b. `tar xzf kafka_2.12-2.2.1.tgz`
- c. `mv kafka_2.12-2.2.1 kafka`

## 2. Create a directory for jmx exporter inside kafka folder

- a. `mkdir prometheus`
- b. `cd prometheus`
- c. `wget https://repo1.maven.org/maven2/io/prometheus/jmx/jmx\_prometheus\_javaagent/0.3.1/jmx\_prometheus\_javaagent-0.3.1.jar`
- d. `wget https://raw.githubusercontent.com/prometheus/jmx\_exporter/master/example\_configs/kafka-2\_0\_0.yml`
- e. Edit the config file of step d by adding the following text at the end (be careful with indentation)
  - pattern: kafka.producer<type=producer-metrics, client-id=(.)><>(.)\w\*
  - name: kafka\_producer\_\$2
  - pattern: kafka.consumer<type=consumer-metrics, client-id=(.)><>(.)\w\*
  - name: kafka\_consumer\_\$2
  - pattern: kafka.consumer<type=consumer-fetch-manager-metrics, client-id=(.)><>(.)\w\*
  - name: kafka\_consumer\_\$2
- f. The config and the exporter java agent will be later on passed to kafka server, kafka producer and kafka consumer to export their metrics (specified above) to prometheus



# Jmx Exporter Configuration Model

- The config file consists of a set of rules to collect attributes of java beans
- Each rule has a pattern to match against bean attribute
  - `domain<beanpropertyName1=beanPropertyValue1, beanpropertyName2=beanPropertyValue2, ...><key1, key2, ...>attrName: value`
- As soon as an attribute matched against a rule, rule processing will stop and attribute is exported to prometheus.

# Start Kafka (1/2)

1. Start zookeeper by running this command from the kafka folder
  - a. `bin/zookeeper-server-start.sh config/zookeeper.properties`
2. Start the Kafka server (replace `USER_NAME` with your username)
  - a. Open another terminal and run this command from the kafka folder

```
KAFKA_HEAP_OPTS="-Xmx1000M -Xms1000M" KAFKA_OPTS="-  
javaagent:/home/USER_NAME/summerSchool/kafka/prometheus/jmx_prometheus_javaagent-  
0.3.1.jar=7071:/home/USER_NAME/summerSchool/kafka/prometheus/kafka-2_0_0.yml" bin/kafka-server-start.sh  
config/server.properties
```

Environment variable for memory requirements of the command

Environment variable for jmx exporter java agent and its config file

Command to start the kafka server

Configuration of the kafka server

# Start Kafka (2/2)

## 3. Start Kafka consumer

- a. Open another terminal and run this command from kafka folder

```
KAFKA_OPTS="-javaagent:/home/USER_NAME/summerSchool/kafka/prometheus/jmx_prometheus_javaagent-0.3.1.jar=7072:/home/USER_NAME/summerSchool/kafka/prometheus/kafka-2_0_0.yml"
/home/USER_NAME/summerSchool/kafka/bin/kafka-console-consumer.sh --bootstrap-server 0.0.0.0:9092 --topic test --from-beginning
```

## 3. Start Kafka producer

- a. Open another terminal and run this command from kafka folder

```
KAFKA_OPTS="-javaagent:/home/USER_NAME/summerSchool/kafka/prometheus/jmx_prometheus_javaagent-0.3.1.jar=7073:/home/USER_NAME/summerSchool/kafka/prometheus/kafka-2_0_0.yml"
/home/USER_NAME/summerSchool/kafka/bin/kafka-console-producer.sh --broker-list 0.0.0.0:9092 --topic test
```

# Install Prometheus

## 1. From the summerSchool folder run

- a. `wget https://github.com/prometheus/prometheus/releases/download/v2.12.0-rc.0/prometheus-2.12.0-rc.0.linux-amd64.tar.gz`
- b. `tar xf prometheus-2.12.0-rc.0.linux-amd64.tar.gz`
- c. `mv prometheus-2.12.0-rc.0.linux-amd64 prometheus`
- d. `cd prometheus/`
- e. Edit `prometheus.yml` by adding this text to the end (specifying the targets that prometheus will get the metrics from) (be careful with indentation)

```
- job_name: 'kafka-server'
  static_configs:
    - targets: ['127.0.0.1:7071']
- job_name: 'kafka-consumer'
  static_configs:
    - targets: ['127.0.0.1:7072']
- job_name: 'kafka-producer'
  static_configs:
    - targets: ['127.0.0.1:7073']
```

# Run the Prometheus server

1. Now run prometheus server with
  - a. `./prometheus --config.file=prometheus.yml`
2. At this point communication of kafka producer and consumer can be tracked from <http://localhost:9090/graph> by choosing metrics with the kafka prefix, for example: `kafka_consumer_incoming_byte_rate`
3. However, a better way is to use grafana for monitoring multiple metrics in a real-time dashboard

☐ Enable query history


kafka\_consumer\_bytes\_consumed\_total

Load time: 182ms  
Resolution: 14s  
Total time series: 1

Execute - insert metric at cursor - ▾

Graph Console

- 1h + ⏪ Until ⏩ Res. (s) ☐ stacked



✓ kafka\_consumer\_bytes\_consumed\_total[instance="127.0.0.1:7072",job="kafka-consumer"]

Remove Graph

Add Graph

# Install & run Grafana

In the previous steps we were able to monitor *individual* metrics of kafka from Prometheus manually. However, Grafana can perform same operation by reading from a config file.

## 1. Create a folder in the summerSchool a folder named Grafana

- a. `mkdir grafana`
- b. `wget https://dl.grafana.com/oss/release/grafana-6.3.3.linux-amd64.tar.gz`
- c. `tar xf grafana-6.3.3.linux-amd64.tar.gz`
- d. `mv grafana-6.3.3 grafana`
- e. `cd grafana/`

## 2. Run Grafana server

- a. `bin/grafana-server start`
- b. Go to <http://localhost:3000> and see the grafana dashboard (user=admin , password=admin)
- c. Add prometheus as a datastore for grafana ("datasource" -> prometheus)
- d. Download the predesigned garfana dashboard for kafka from here and unzip to get the json file  
[https://activewizards.com/content/blog/Kafka\\_Monitoring\\_with\\_Prometheus\\_Telegraf\\_and\\_Grafana/grafana-system-and-kafka.zip](https://activewizards.com/content/blog/Kafka_Monitoring_with_Prometheus_Telegraf_and_Grafana/grafana-system-and-kafka.zip)
- e. Import dashboard json in grafana dashboard by going to Dashboard > home > + > import json

# Monitor Kafka message exchange metrics from Grafana

By exchanging messages on kafka between producer and consumer, kafka metrics will change and can be observed from the Grafana dashboard in real-time.



# Prometheus Rules

- You can extend Prometheus with some rules for monitoring and alerting
- This extension is applied to prometheus server after adding following text to the prometheus config file (prometheus/prometheus.yml)

```
rule_files:  
  - 'prometheus.rules.yml'
```

- Create a new file in this directory named Prometheus.rules.yml
- Add the following alert definition to Prometheus.rules.yml file

```
groups:  
  - name: example  
    rules:  
  
    # Alert when kafka consumer consumes more than 100 bytes  
    - alert: ConsumerOverload  
      expr: kafka_consumer_bytes_consumed_total > 100  
      for: 2s  
      labels:  
        severity: page  
      annotations:  
        summary: "consumer instance {{ $labels.instance }} overloaded"
```

# Homework

- Try to train machine learning models for various metrics
- Currently Prometheus expr only accepts mathematical expression
- Extension to prometheus for accepting ML models in the expr field

```
groups:  
- name: example  
  rules:  
  
    # Alert when kafka consumer consumes more than 100 bytes  
    - alert: ConsumerOverload  
      expr: kafka_consumer_bytes_consumed_total > 100  
      for: 2s  
      labels:  
        severity: page  
      annotations:  
        summary: "consumer instance {{ $labels.instance }} overloaded"
```