

# Software evolution in microservice artefacts

Josef Spillner <josef.spillner@zhaw.ch>  
Service Prototyping Lab ([blog.zhaw.ch/splab](http://blog.zhaw.ch/splab))

Sep 03, 2019 | INFORTE Summer School, Tampere

# I. Artefacts Primer



# Software-as-a-Service

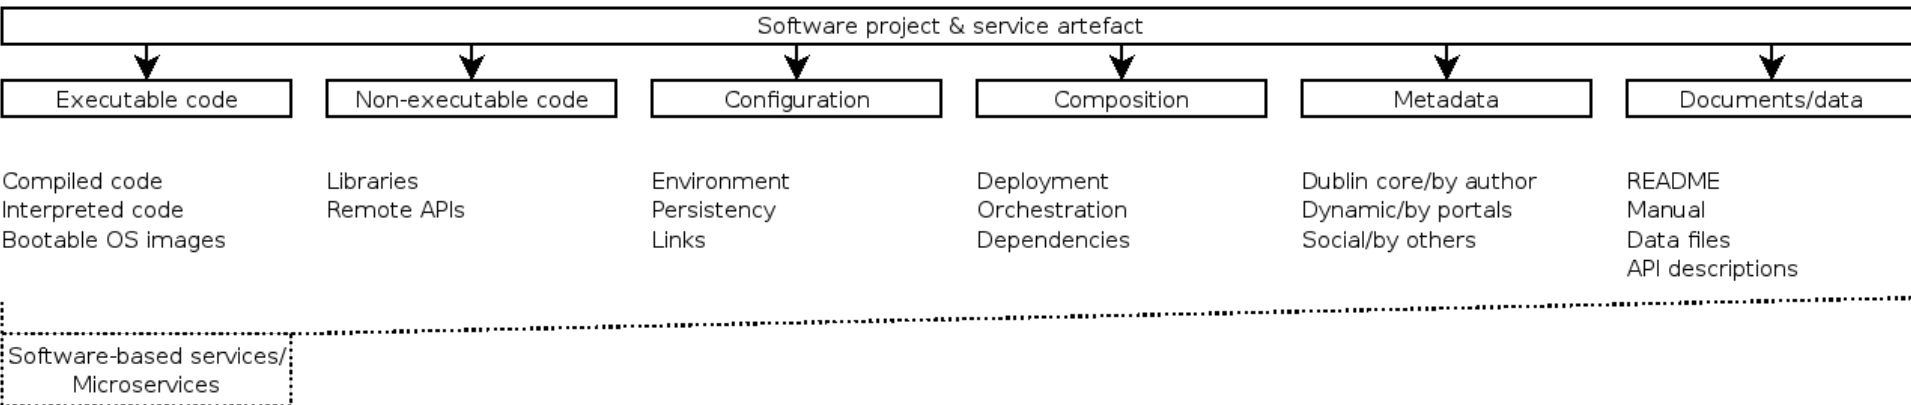
## (Micro-)Service

- uniform description and invocation
  - description languages
  - messaging protocols & exchange patterns
- discoverable
  - registries
- encapsulation of implementation / separation of interface
  - software implementation is common (often polyglot)
    - SaaS, also: PaaS, FaaS, ... (→ subset of XaaS)
- composable/non-discriminating
- networked
- access models (free, subscription, contract, ...)



# Software artefacts

- composite software - mixed-technology artefacts
- multiple file types often combined in artefact technologies
- shared via repositories/hubs/marketplaces



(often: multi-category artefacts, e.g. deployment for SaaS = composition + configuration + metadata)



# Software code & non-code artefacts

Software applications/services := bundle of artefacts

- typical formats: WAR, SAR, BPR, ZIP, ..., Helm charts

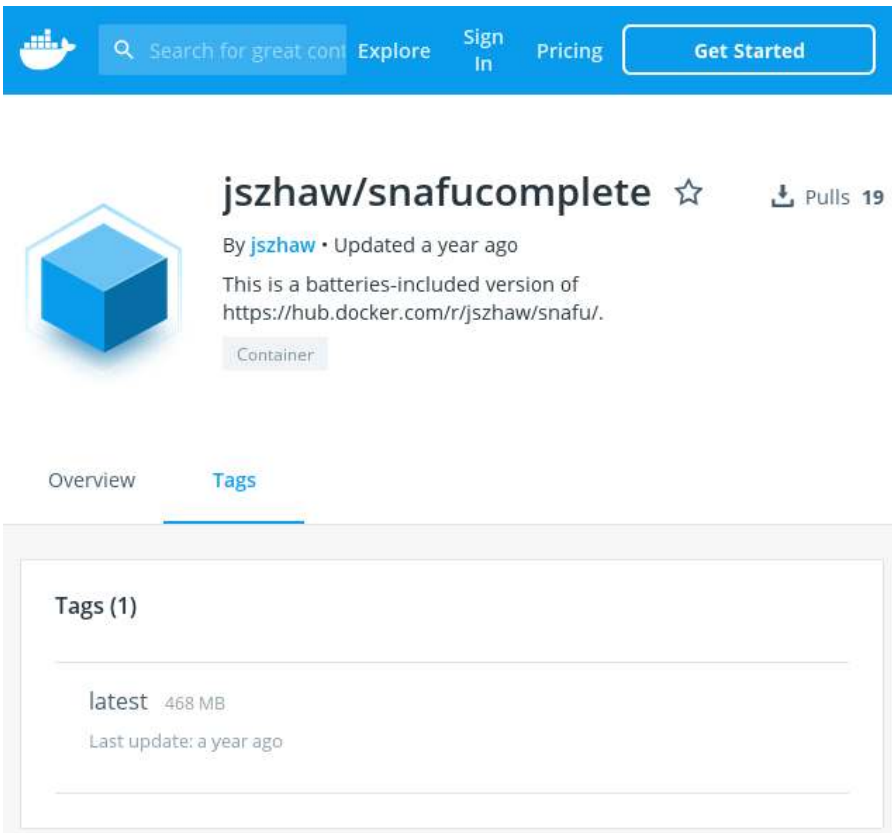
Artefact categories

- executable and non-executable code
  - \*.class, \*.jar, \*.py(c), ...
- service descriptions
  - WSDL, WADL, WSML, GUIDD, USDL, RAML, Swagger/OpenAPI
- SLA templates
  - WSLA, WS-Agreement
- deployment descriptors
  - infrastructure: HOT, Cloudformation etc.
  - BPEL: PDD
  - servlets: web.xml
  - Kubernetes: YAML/JSON files, Helm templates
  - Lambda functions: SAM templates
  - ...



# Metadata

## Example: Docker Hub



The screenshot shows the Docker Hub interface for the repository `jszhaw/snafucomplete`. The repository is marked as a container and has 19 pulls. The description states it is a batteries-included version of `https://hub.docker.com/r/jszhaw/snafu/`. The 'Tags' tab is selected, showing a single tag `latest` with a size of 468 MB and a last update of a year ago.



```
1 {  
2   "name": "image_name",  
3   "data": {  
4     "1560176146": {  
5       "pull_count": 1721041870,  
6       "star_count": 6989,  
7       "is_migrated": false,  
8       "is_automated": false,  
9       "number_of_tags": 278,  
10      "full_size": 35224772,  
11      "os": "linux",  
12      "architecture": "amd64",  
13      "latest": true  
14    }  
15  }  
16 }
```



→ access: API (results limit), web crawling (hits frequency limit)

# Metadata

Typical content:

- “Dublin Core metadata” for authored works
  - title, creator, subject, description, publisher, contributor, ..., rights
- developers/maintainers with contact information
- version specification
- dependencies
- social: ratings, comments
- auto-generated: downloads

Example: Helm’s Chart.yml

dependencies:

inferred from present  
subdirectories charts/\*

```
apiVersion: v1
name: ExampleChart
version: 1.0
kubeVersion: 1.7a
description: This is just an example chart.
keywords:
  - example
  - nonsense
home: http://inforte.jyu.fi/events/SW_evolution
sources:
  - https://github.com/serviceprototypinglab
maintainers:
  - name: Josef Spillner
    email: josef.spillner@zhaw.ch
    url: https://blog.zhaw.ch/splab/
engine: gotpl
icon: icons/myicon.png
appVersion: 2.0alpha2
deprecated: true
tillerVersion: The version >2.0.0
```

# Code

## Programming language

- Java, Python, Go, JavaScript, Solidity, ...

## Language version and (runtime) flavour

- Java 6 vs. Java 9, CPython vs. PyPy

## Representation

- source code, byte code, compiled code
- flat file, archive file, locator
- include paths

## Context

- completeness, up-to-dateness, genuineness
- degree of deviation
- automated build instructions
- relationships to other code units

```
// Global (instance-wide) scope; this
// computation runs at instance cold-start
const instanceVar = heavyComputation();

/**
 * HTTP Cloud Function declaring a variable.
 *
 * @param {Object} req
 *   Cloud Function request context.
 * @param {Object} res
 *   Cloud Function response context.
 */
exports.scopeDemo = (req, res) => {
  // Per-function scope; this
  // computation runs every time
  // this function is called
  const functionVar = lightComputation();

  res.send(`Per instance: ${instanceVar},
    per function: ${functionVar}`);
};
```



private cloud





# Declarative deployment descriptors

Deployment: code (scripts) vs. declarative descriptor

Typical content:

- name within the target system
- runtime limits
- access control
- logging setup
- network setup: port numbers, throttling, ...

Example: Helm's deployment template for Joomla  
(not fully declarative;  
some metadata present)

```
{{- if or .Values.mariadb.enabled .... -}}
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: {{ template "joomla.fullname" . }}
  labels:
    app: {{ template "joomla.fullname" . }}
    chart: {{ template "joomla.chart" . }}
    release: {{ .Release.Name | quote }}
    heritage: {{ .Release.Service | quote }}
spec:
  selector:
    matchLabels:
      app: {{ template "joomla.fullname" . }}
      release: "{{ .Release.Name }}"
  replicas: 1
  template:
    metadata:
      labels:
        app: {{ template "joomla.fullname" . }}
    spec:
      {{- if .Values.image.pullSecrets }}
      imagePullSecrets:
        {{- range .Values.image.pullSecrets }}
        - name: {{ . }}
      {{- end}}
      {{- end}}
      containers:
```

.....

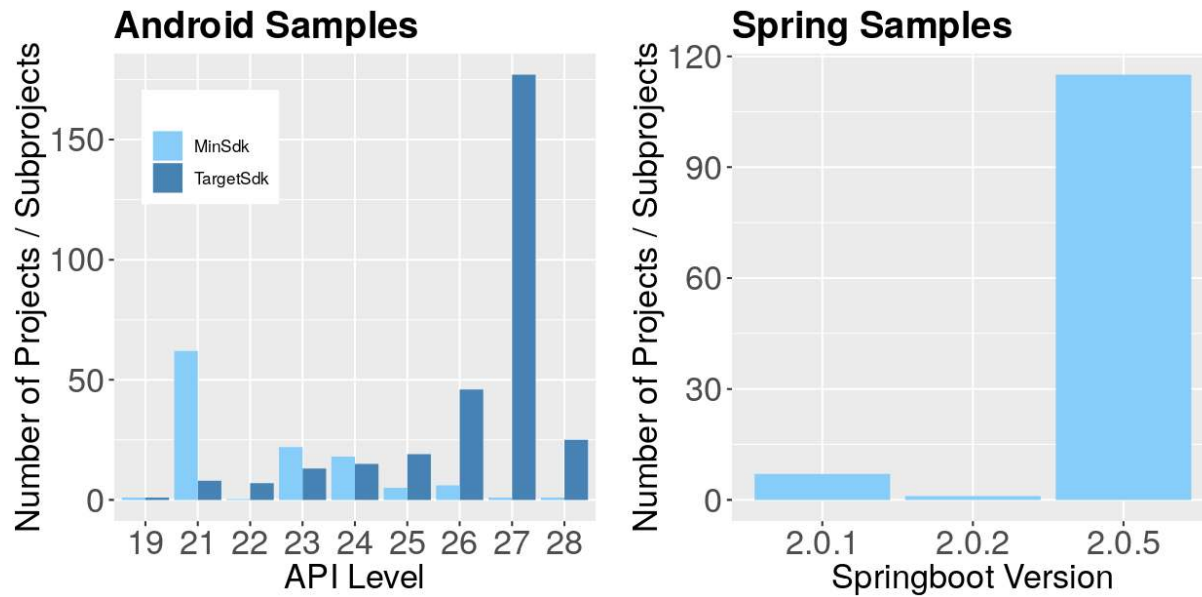


# Software artefacts evolution

Implicit hypothesis: Correlation with software evolution

## Lifecycle

- branches & releases
- numbering schemes
- change patterns over time and/or over versions

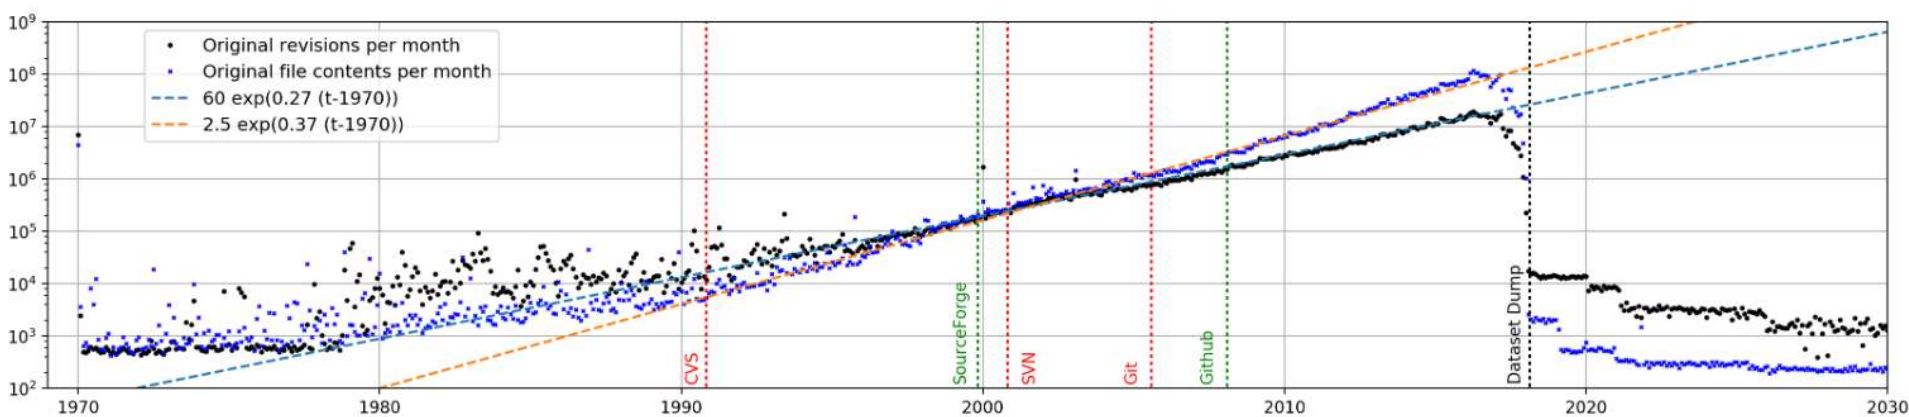


G. Menezes, B. Cafeo, A. Hora: Framework Code Samples: How Are They Maintained and Used by Developers? arXiv:1907.05564

# Code vs. artefact considerations

Software Heritage study - 4 billion code file artefacts, 1 billion commits

- exponential growth rate
- missing link to artefacts, but strong link to repository technologies
  - higher-level features & release/duplication efforts
- data point: since SourceForge inception,
  - original revisions double every ~30mo
  - original code files double every ~22mo
- data point: 70% of code on GitHub → file-level exact clones

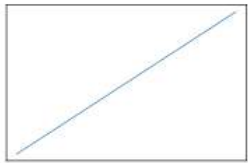


G. Rousseau, R. Di Cosmo, S. Zacchioli: Growth and Duplication of Public Source Code over Time: Provenance Tracking at Scale.  
arXiv:1906.08076

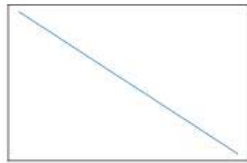
# Software evolution patterns

(per metric)

(unsuitable names)



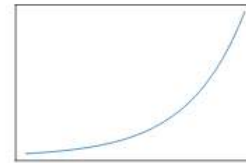
(a) Constant  
Rise



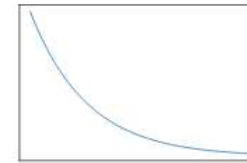
(b) Constant  
Decline



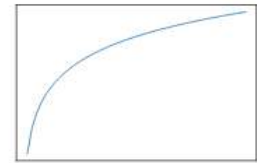
(c) Stability



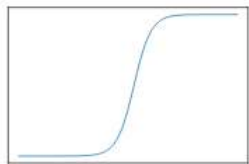
(d) Sudden  
Rise



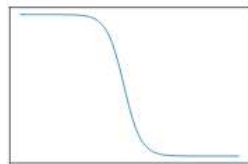
(e) Sudden  
Decline



(f) Sudden Rise  
Plateau



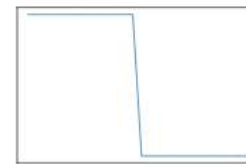
(g) Plateau  
Gradual Rise



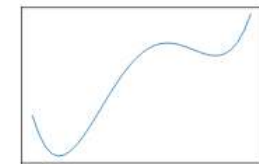
(h) Plateau  
Gradual  
Decline



(i) Plateau  
Sudden  
Rise



(j) Plateau  
Sudden  
Decline



(k) Instability



# II. Practical Examples



# Excursus: Docker Compose files

Compose: Tool for defining and running multi-container applications.

- services, networks, volumes

Version 1 - Nov 2015

Version 2 - Feb 2016

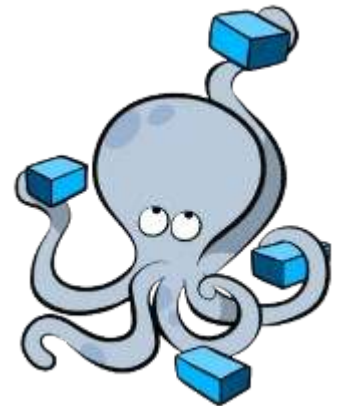
Version 3 - Jan 2017

Quality checks:

containers (external tool): *hadolint*, compositions: *none*

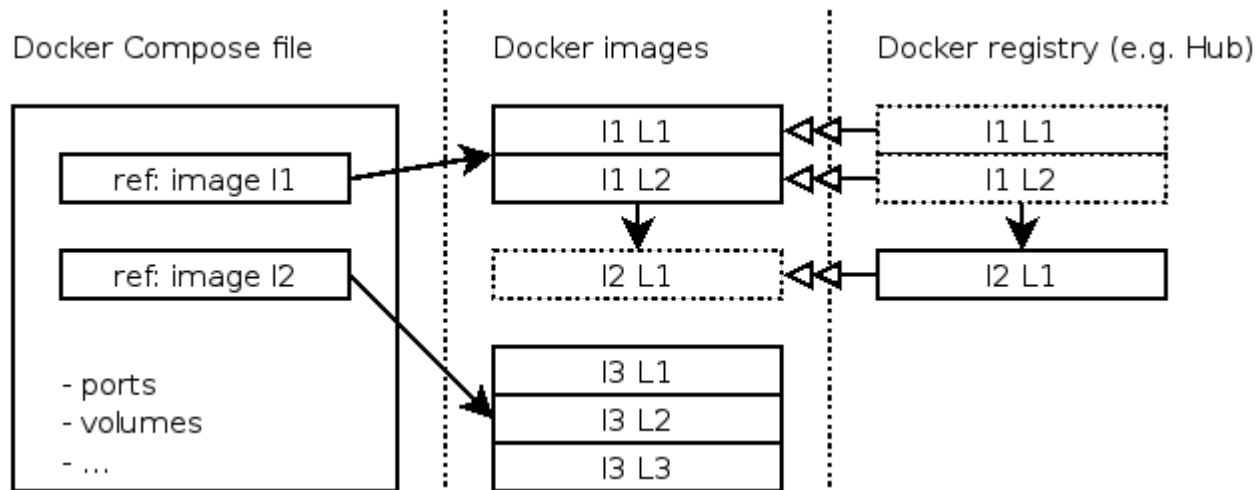
Repositories:

only generic code hosting, no dedicated hub site



# Docker Compose in detail

## Single compose file



# Docker Compose artefact quality

## Duplication

- common in copy&paste-style development
- duplication vs. unique assignments
  - service name
  - container name
  - image
  - port

## Invalid specifications

- common when changing environments, e.g. dev-prod or dev1-dev2
- path references and remote locators
  - volume directory
  - image/version in non-standard registry

```
1  version: '3'
2  services:
3    nginx:
4      container_name: some-nginx
5      image: nginx:1.13
6      restart: always
7      ports:
8        X 3306:80
9        - 443:443
10     expose:
11       - '80'
12     volumes:
13       - ./nginx/conf.d:/etc/nginx/conf.d
14     depends_on:
15       - app
16
17    nginx:
18      container_name: some-mysql
19      image: mysql/mysql-server:5.7
20      environment:
21        MYSQL_DATABASE: test
22        MYSQL_ROOT_PASSWORD: hellokoding
23        MYSQL_ROOT_HOST: '%'
24      ports:
25        X "3306:3306"
26      restart: always
```





# Excursus: Helm apps & charts

Helm: Package manager for the cloud-native landscape



Created in 2015 as a Deis (now MS) installer

Out of K8s incubator in 2017

80 stable charts, 53 maints

KubeApps UI/marketplace (KubeApps Hub) in 2017

Helm summit in February 2018 with 179 participants (out of 250)

CNCF top-level (incubator) project since May 2018

177 stable charts, 142 maints

50k montly downloads

>30k charts on Github/similar?

Quality checks (v2.8): *helm lint*

»This command takes a path to a chart and runs a series of tests to verify that the chart is well-formed.«

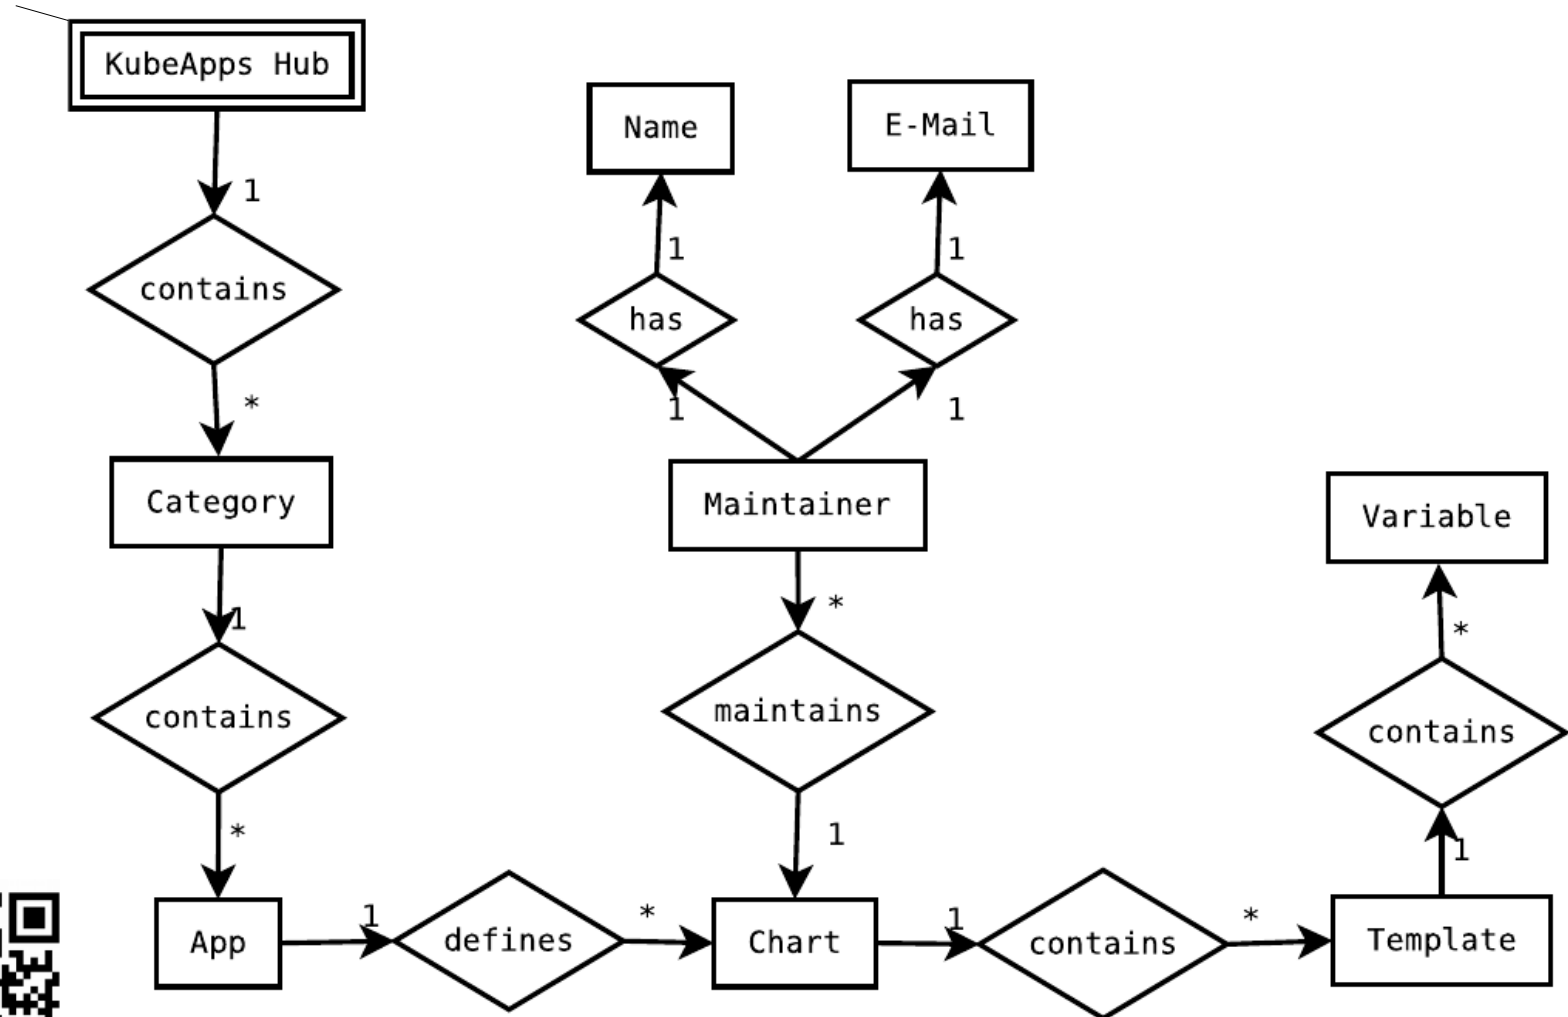
sufficient?

Now: Helm Hub & Vand.io, >700 total charts (all branches),  
ca. 300 stable - e.g.: `helm install stable/drupal` to get CMS



# Helm in detail: entities & relationships

similar: Helm Hub, Vand.io, ...

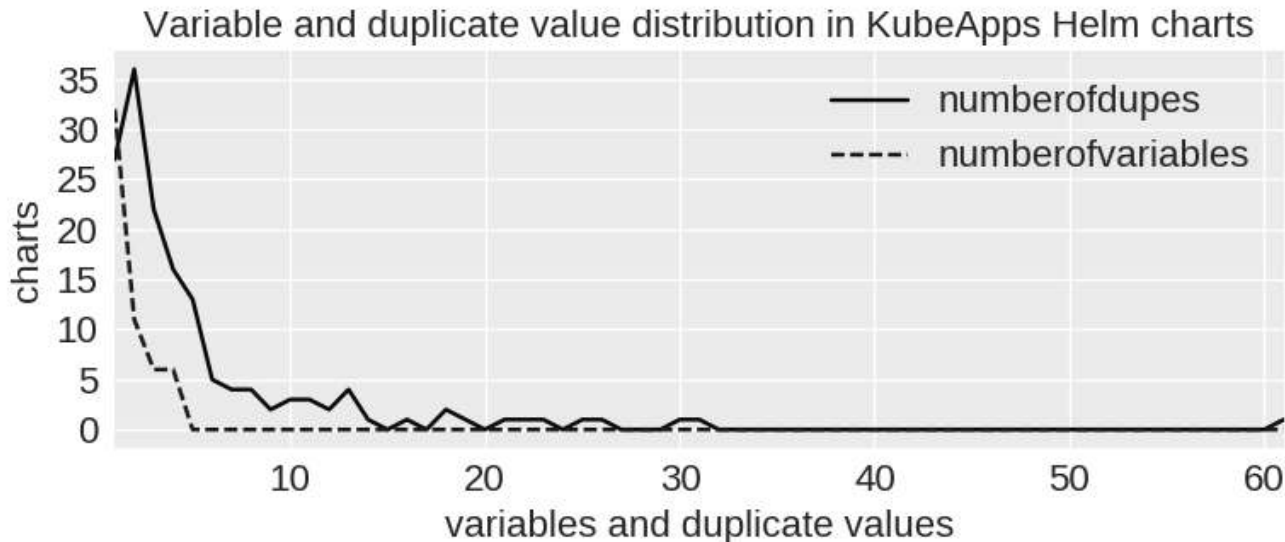


# Helm artefact quality

Equivalence to code smells? Anti-patterns? Non-adherence to best practices?

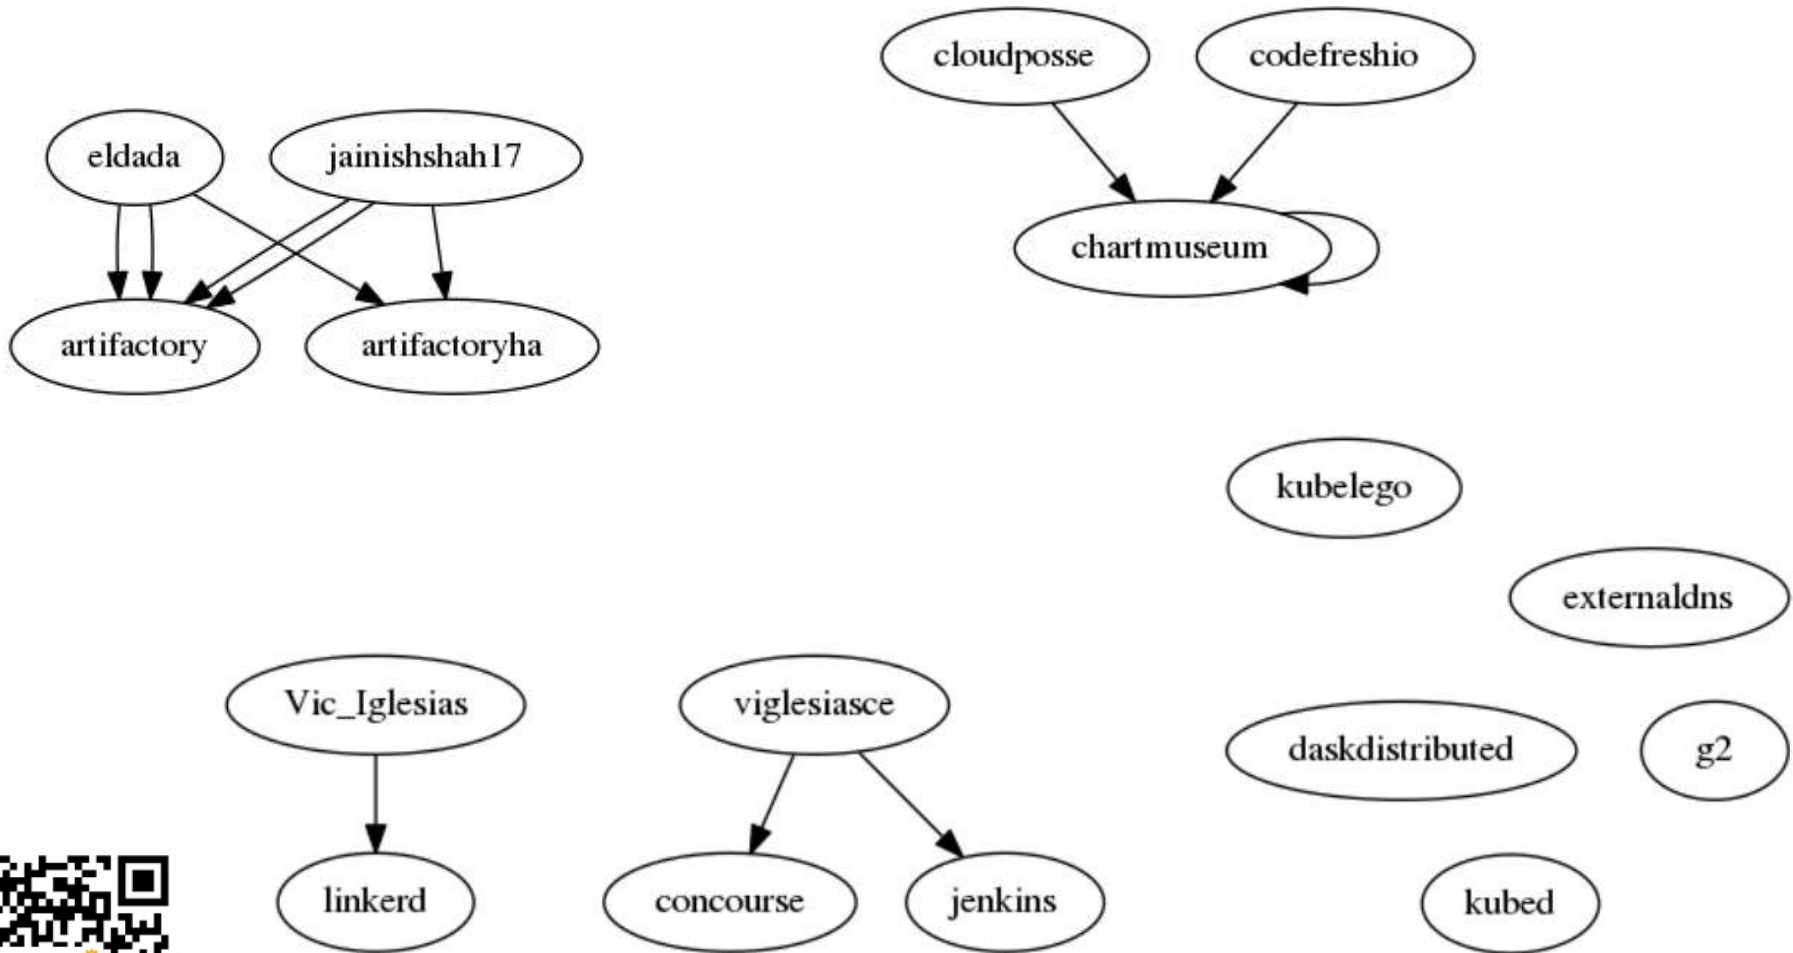
Concrete issues with Helm charts:

- multiple versions in one repository category (e.g., “stable”)
- equivalence of chart name and maintainer name
- duplicate values not expressed as variables
- (dependencies: transitive issues, error propagation & aggravation)



# Helm artefact quality

Invalid maintainer → chart relationship graph (DAG) identification



# Helm artefact quality

Effect of mandatory (albeit low-barrier) quality gates

Chart-related KubeApps Hub metrics

Chart metrics	Affected	Percentage	Base metric
Charts	177	100.00	–
Irregularity: no maintainer	10	5.65	Charts
Irregularity: name collision	1	0.56	Charts
Irregularity: multiple versions	5	2.82	Charts

Chart-related GitHub metrics

Chart metrics	Affected	Percentage	Base metric
Charts	115	100.00	–
Irregularity: no maintainer	95	82.61	Charts
Irregularity: name collision	2	1.74	Charts
Irregularity: multiple versions	0	0.00	Charts

\*without fork analysis



# Excursus: AWS Lambda & SAM

Lambda: Cloud functions running as short-lived containers



AWS Lambda

```
def lambda_handler(event, context):  
    "event: dict  
    context: meta information obj  
    returns: dict, string, int, ..."  
    # ...  
    return "result"
```

SAM: Serverless Application Model

- vendor-specific cloud function bundle technology



## Datadog-Log-Forwarder

Datadog

[https://github.com/DataDog/datadog-serverless-functions/tree/master/aws/logs\\_monitoring](https://github.com/DataDog/datadog-serverless-functions/tree/master/aws/logs_monitoring)

102 stars

Deploy

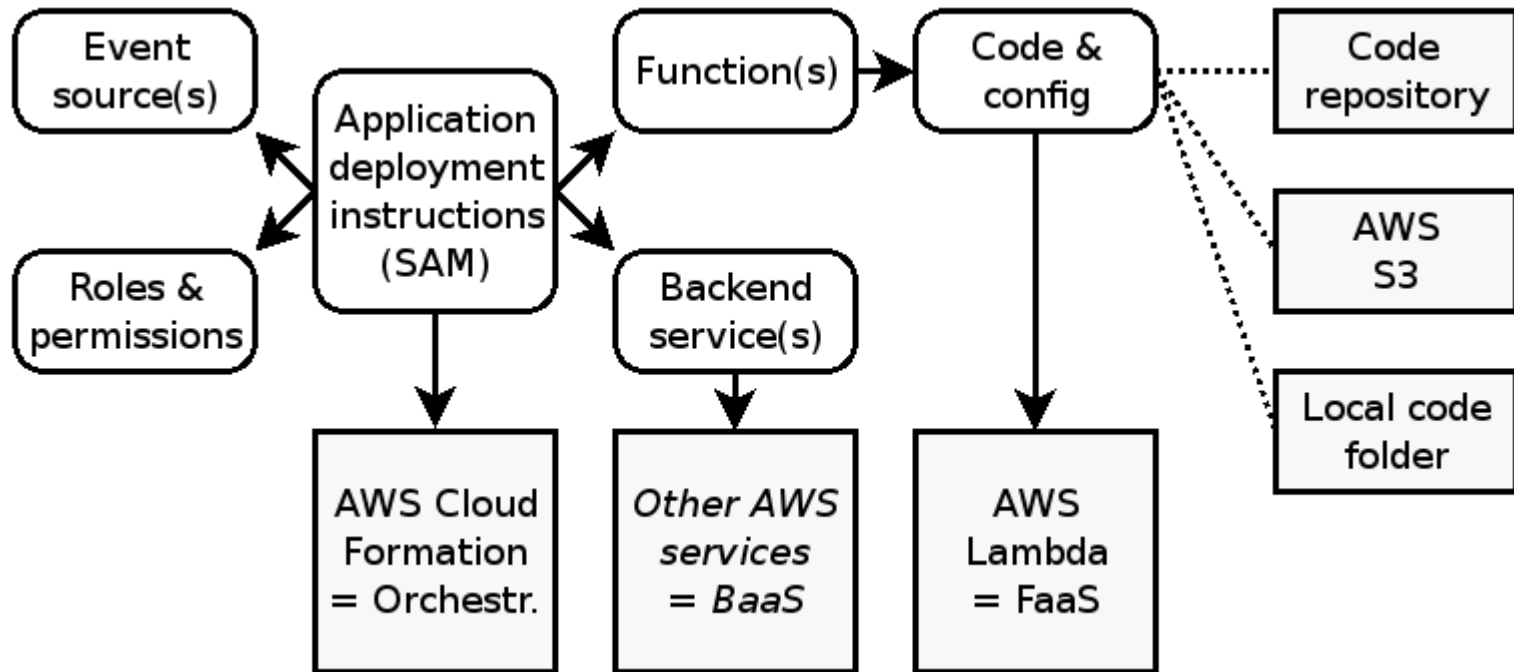
936 deployments

The Datadog log forwarder ships logs stored in S3 and CloudWatch Logs to Datadog for live search, analytics, and alerting.

public repository  
available

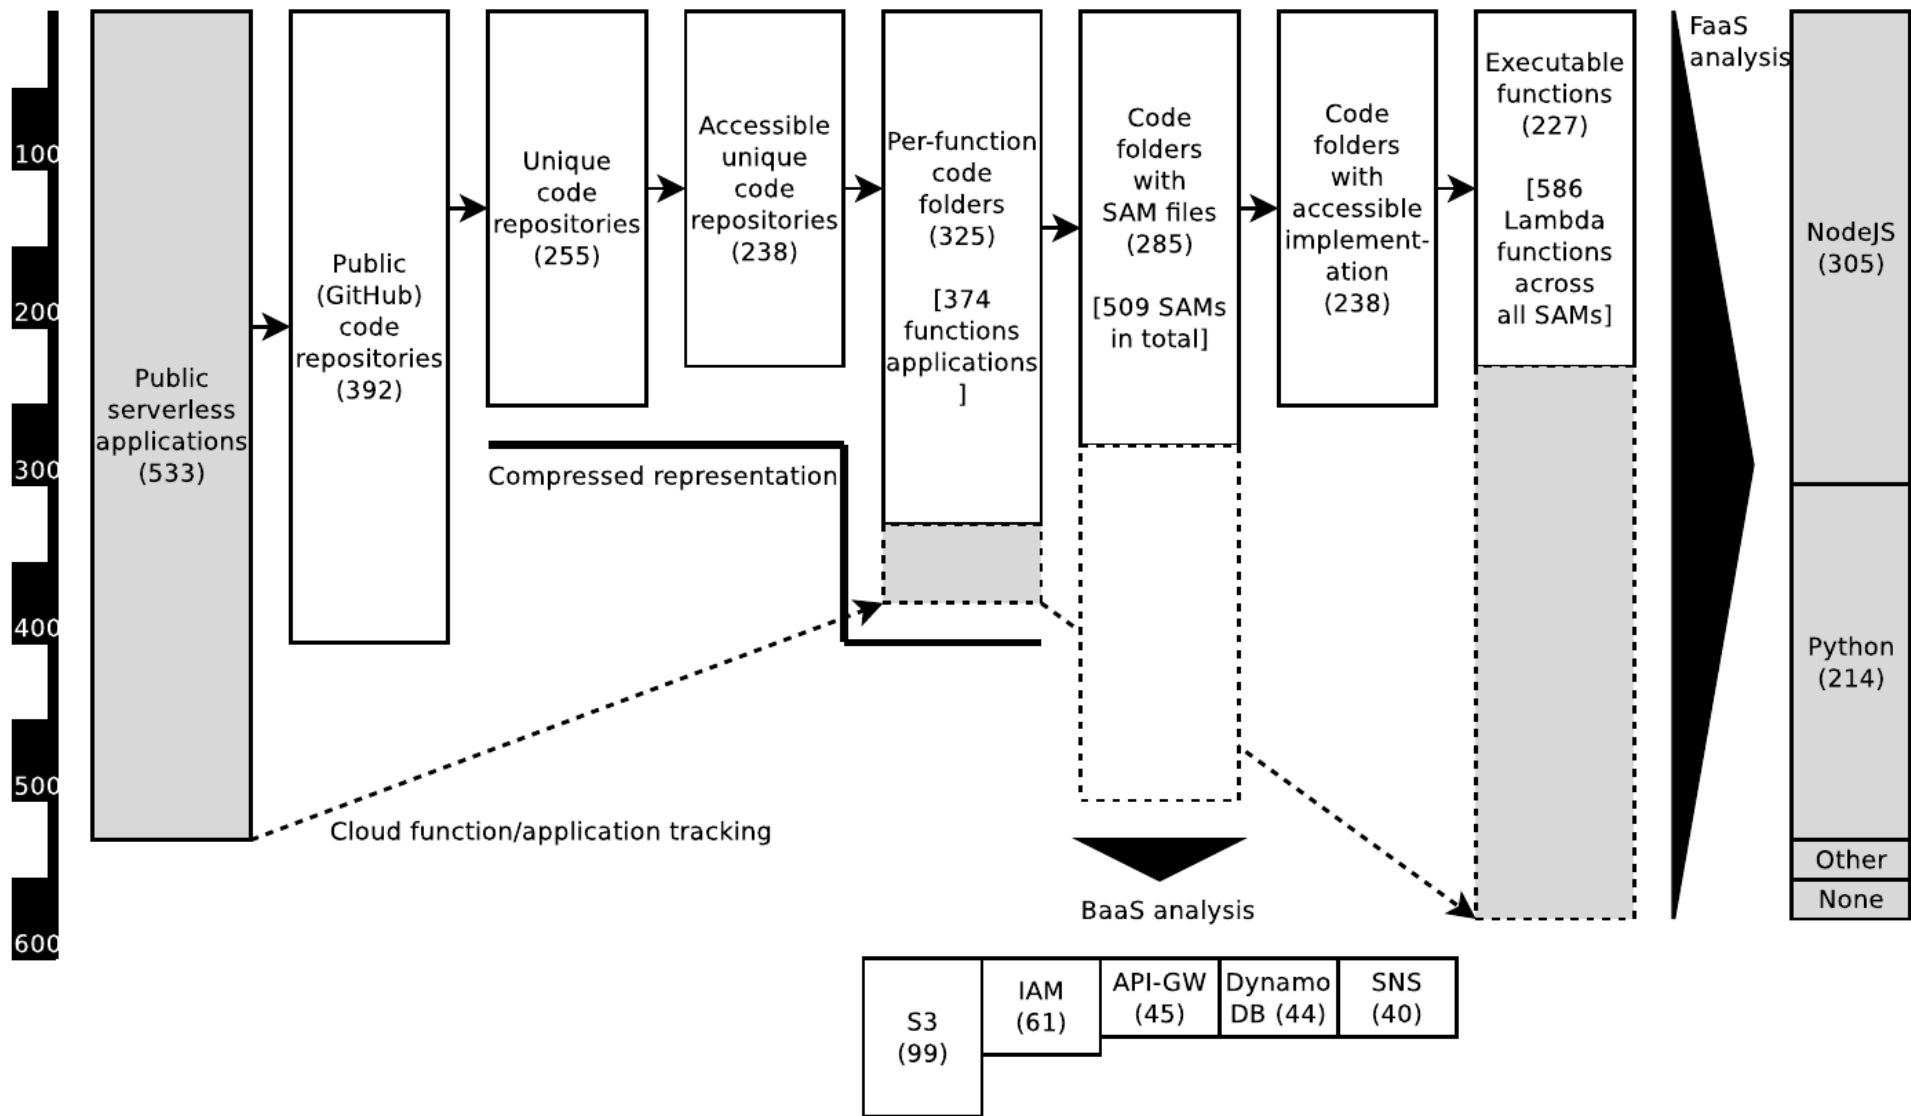
# Lambda/SAM in detail: topology

AWS-defined artefact topology



# Lambda/SAM artefact quality

## Drill-down approach: Issues at SAR level

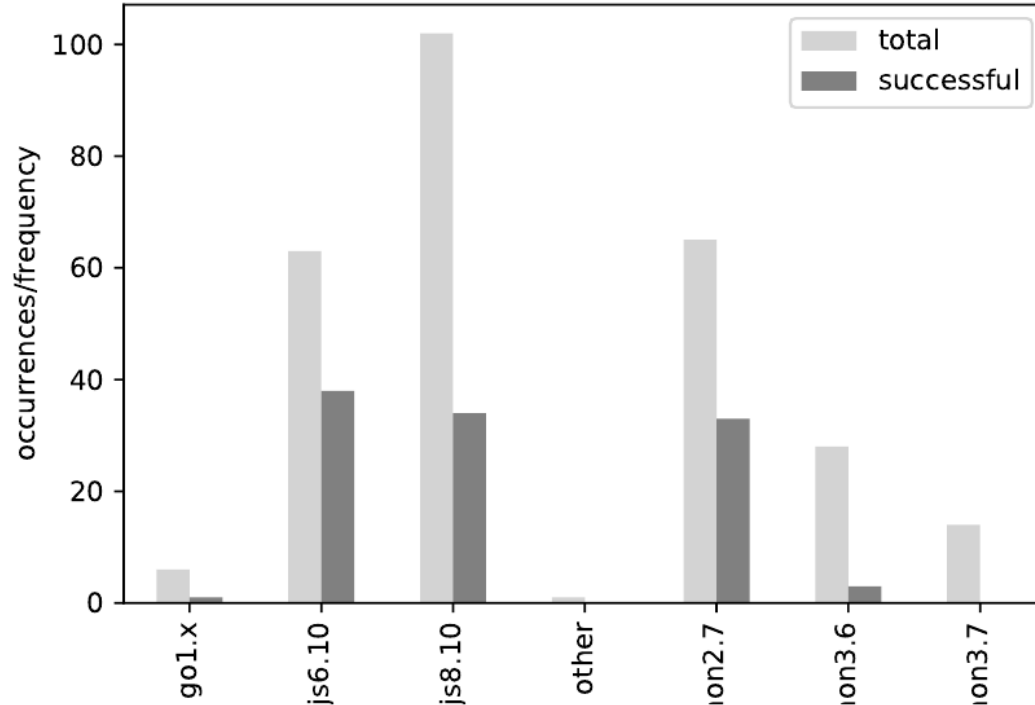




# Lambda/SAM artefact quality

## Generic execution measurements

- caveat: incomplete test event generation + semantics



category	execution time (ms)	billed time (ms)	memory size (MB)	memory used (MB)
successful	78.36	174.23	131.30	22.86
failed	131.31	213.81	295.32	27.18
overall	106.76	195.46	301.92	26.80



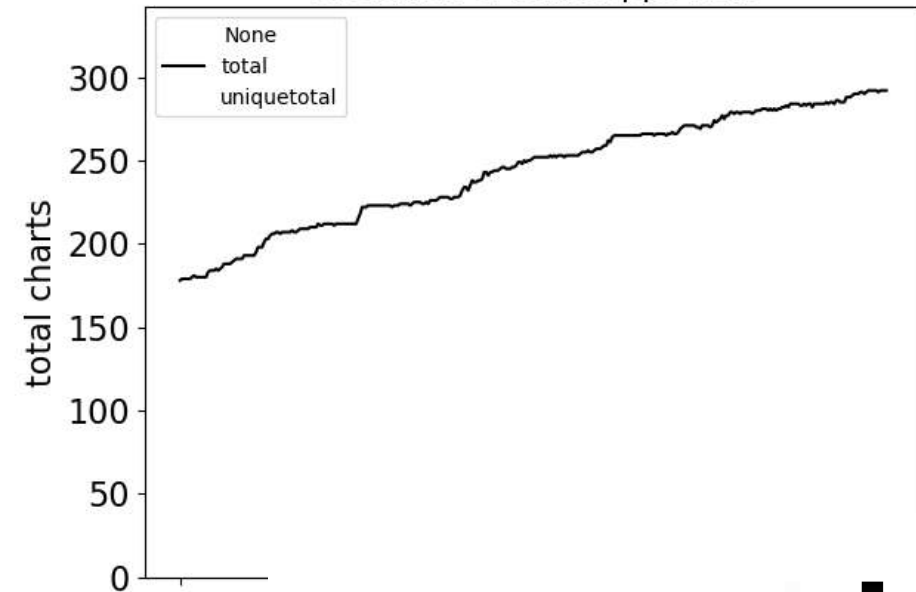
# III. Evolution Observation



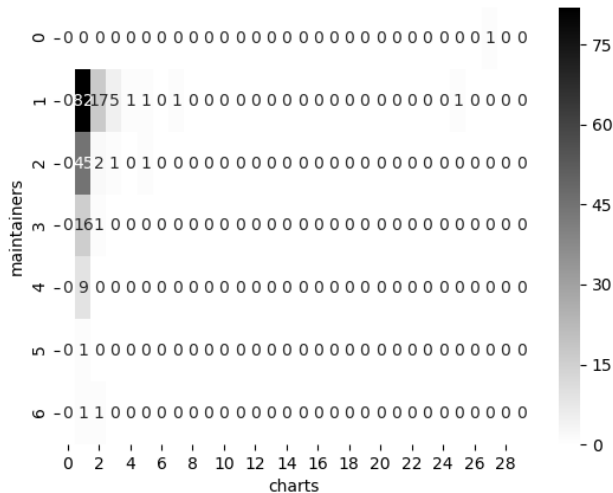
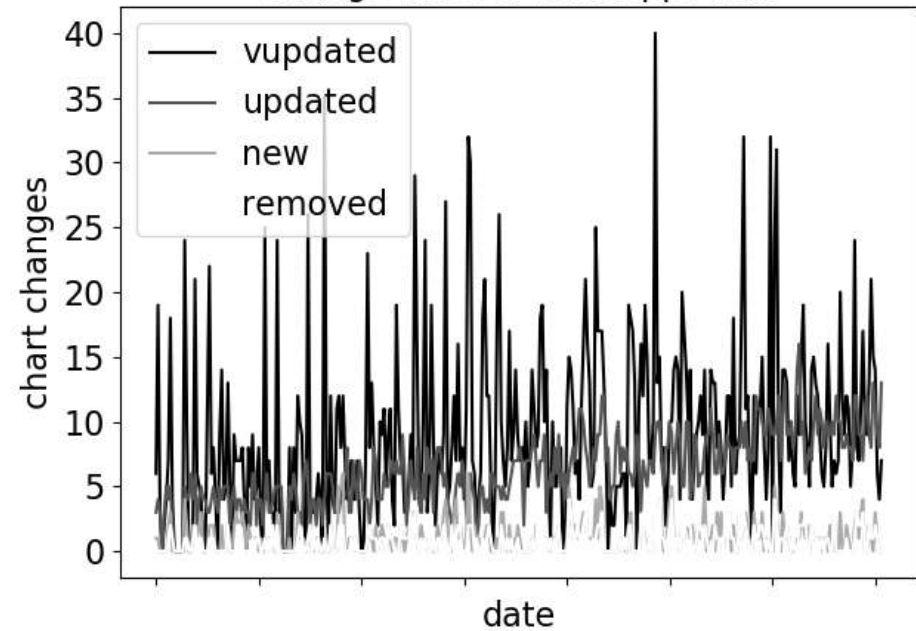
# Helm evolution

Helm charts (May'18-May'19):

Evolution of KubeApps Hub

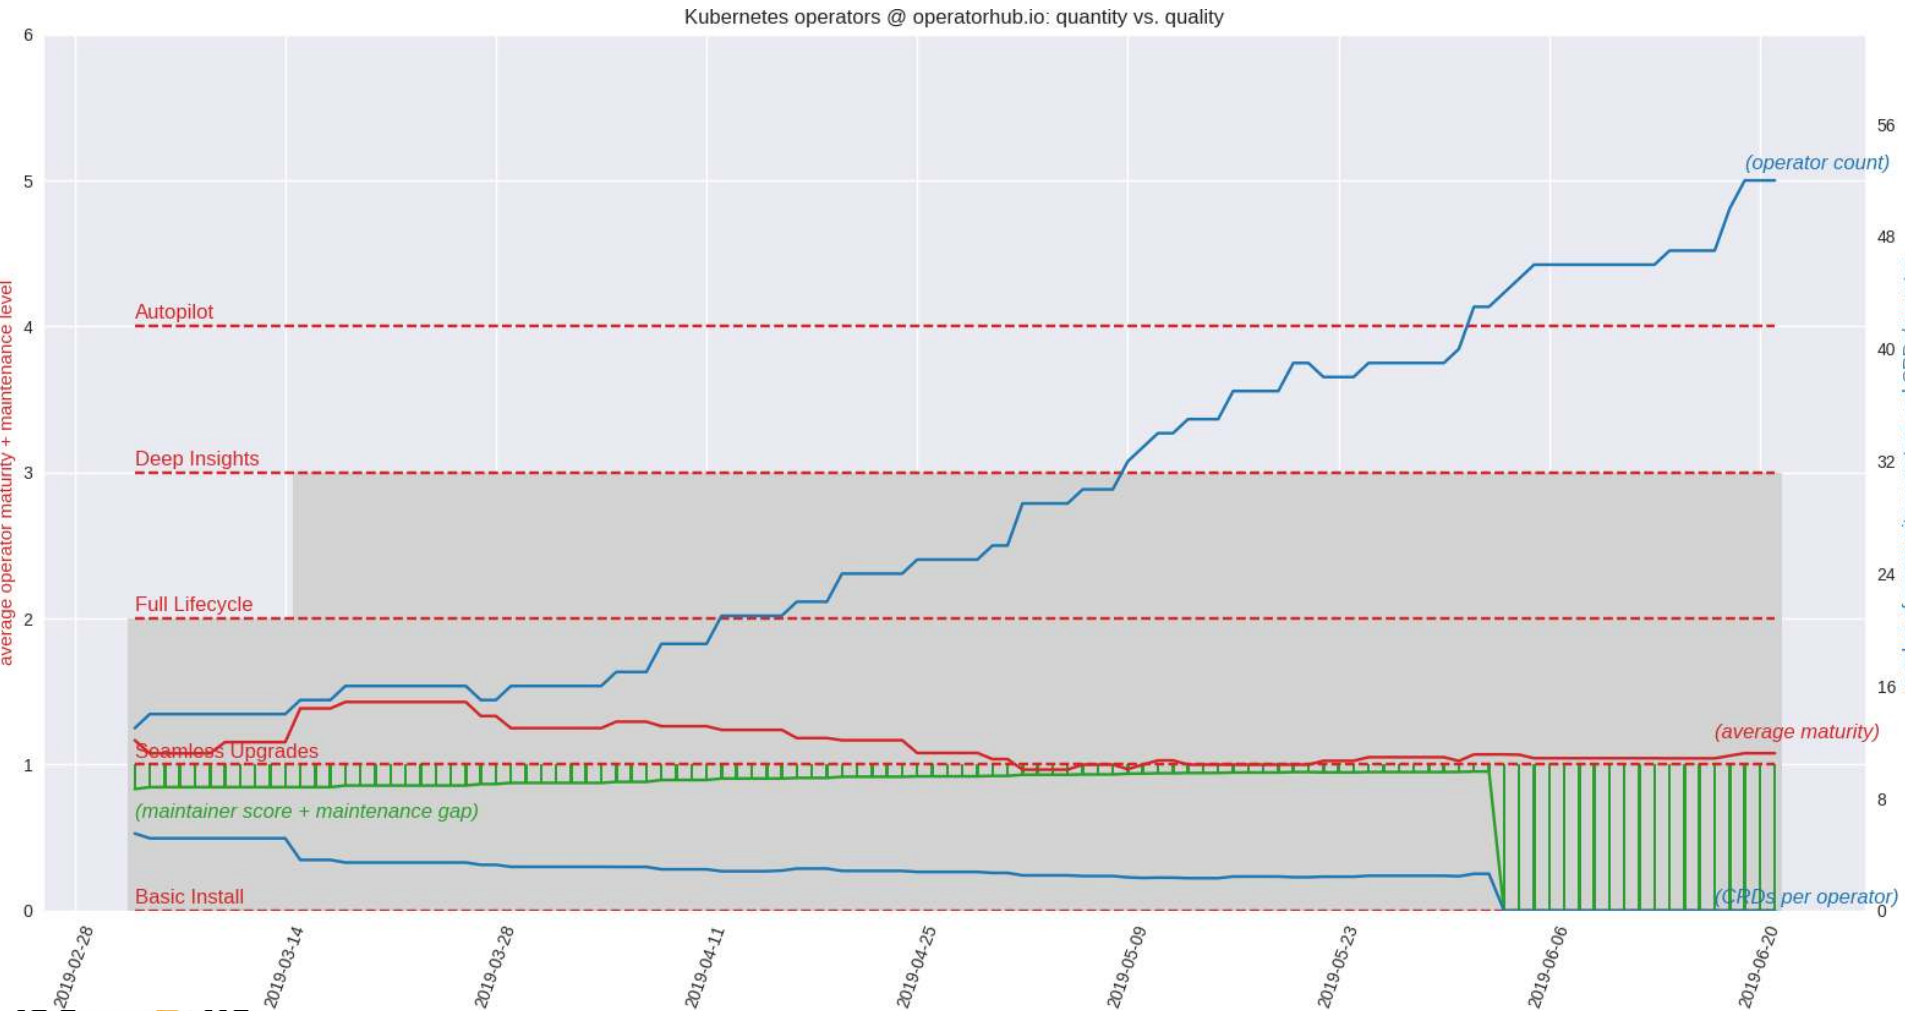


Change rates of KubeApps Hub



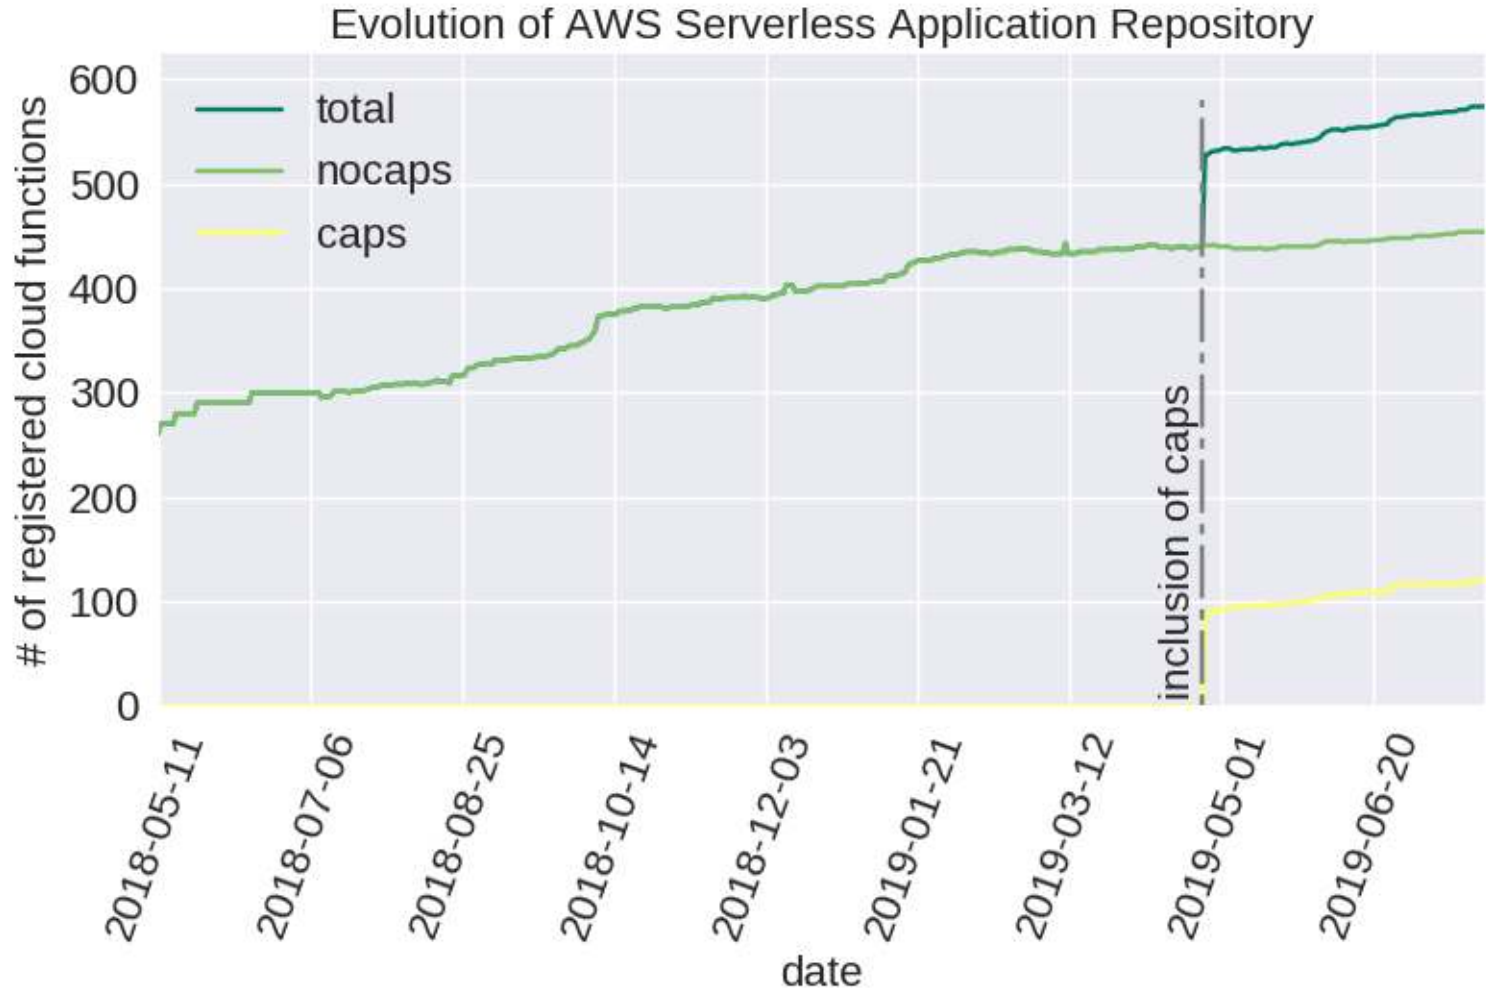
# Kubernetes operators evolution

Operators (Feb'19-Jun'19):



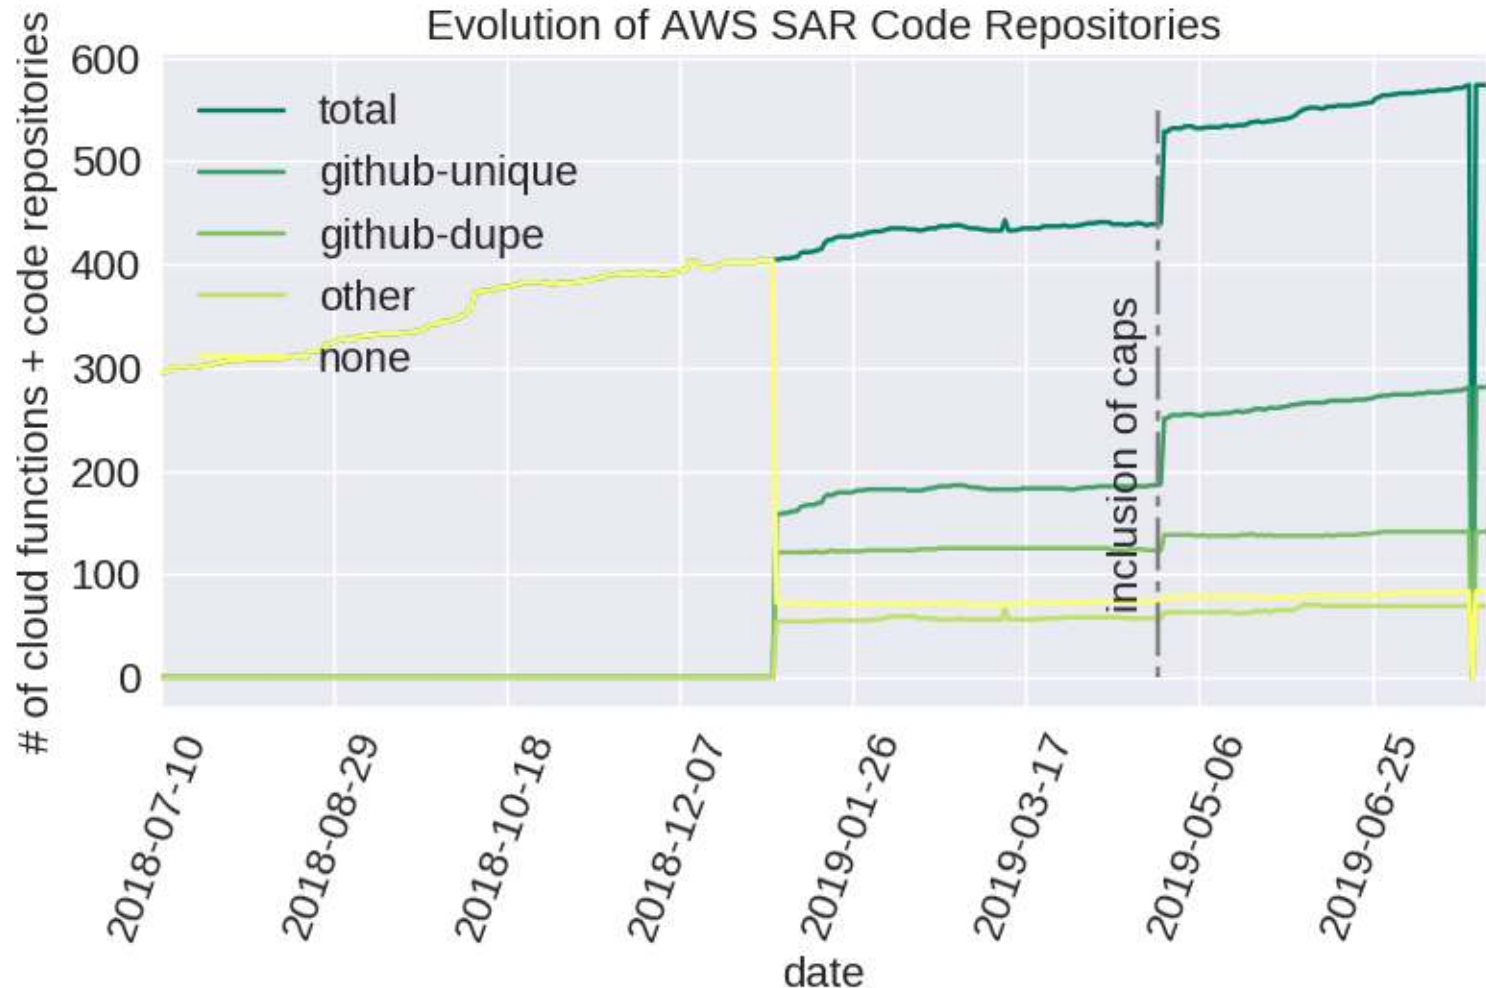
# Lambda/SAM evolution

SAMs at AWS SAR (May'19-Jul'19)



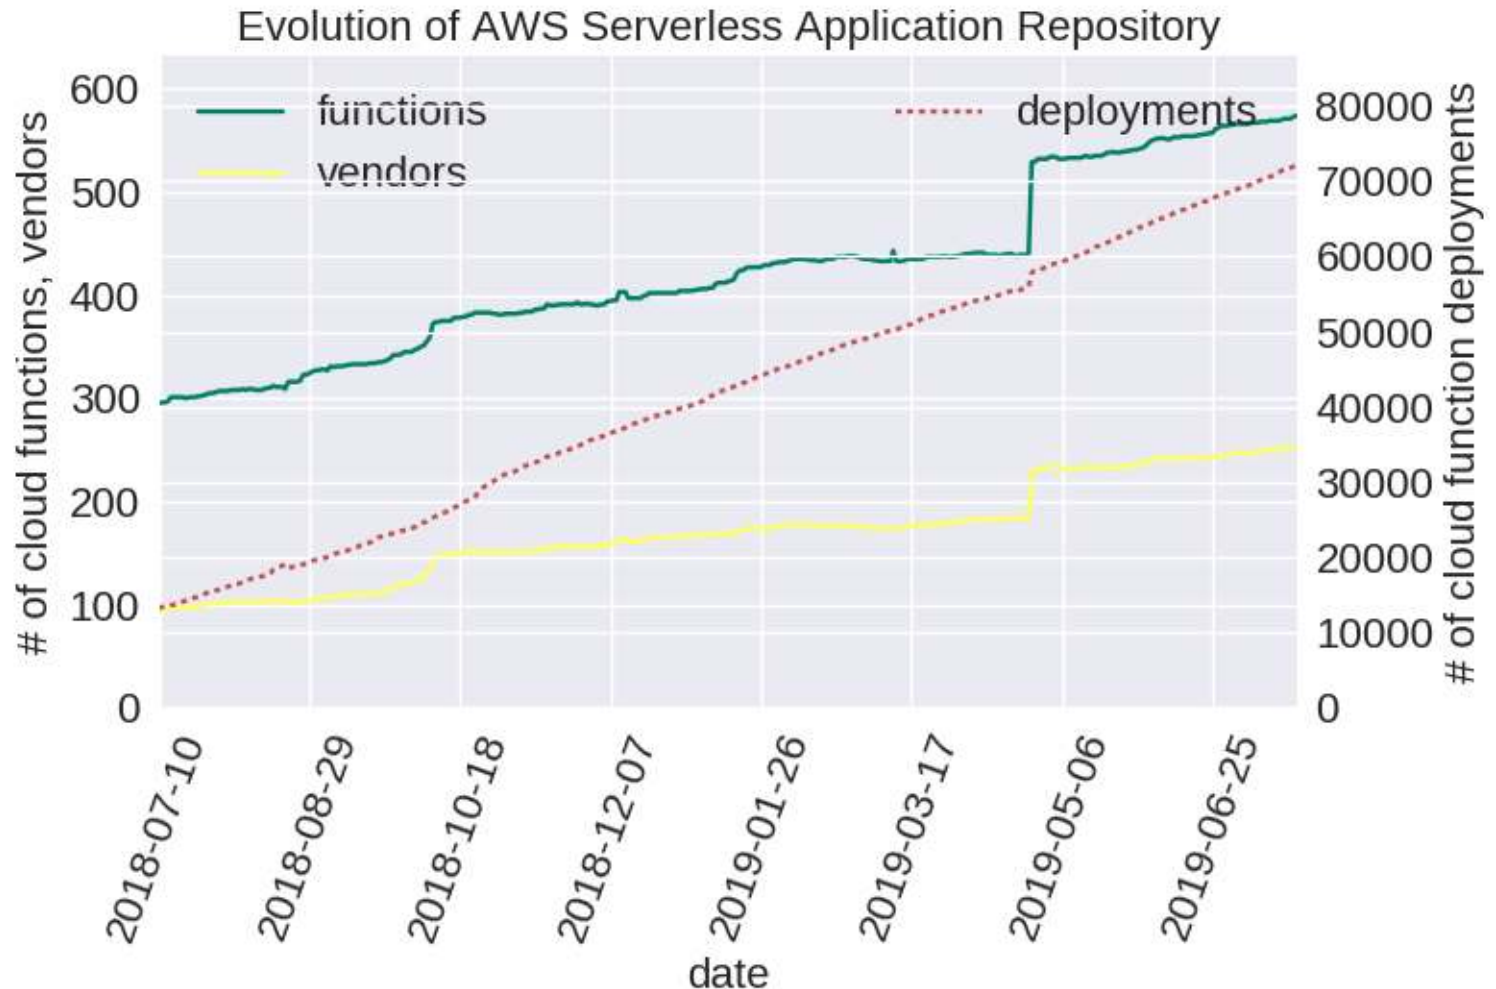
# Lambda/SAM evolution

SAM-referenced code repositories at AWS SAR (Jul'19-Jul'19)



# Lambda/SAM evolution

Vendors and deployments at AWS SAR (Jul'19-Jul'19)



# Remarks on evolution tracking

Developer interest:

- on-demand
- real-time QA (e.g. in CI/CD step)
- for limited set of artefacts

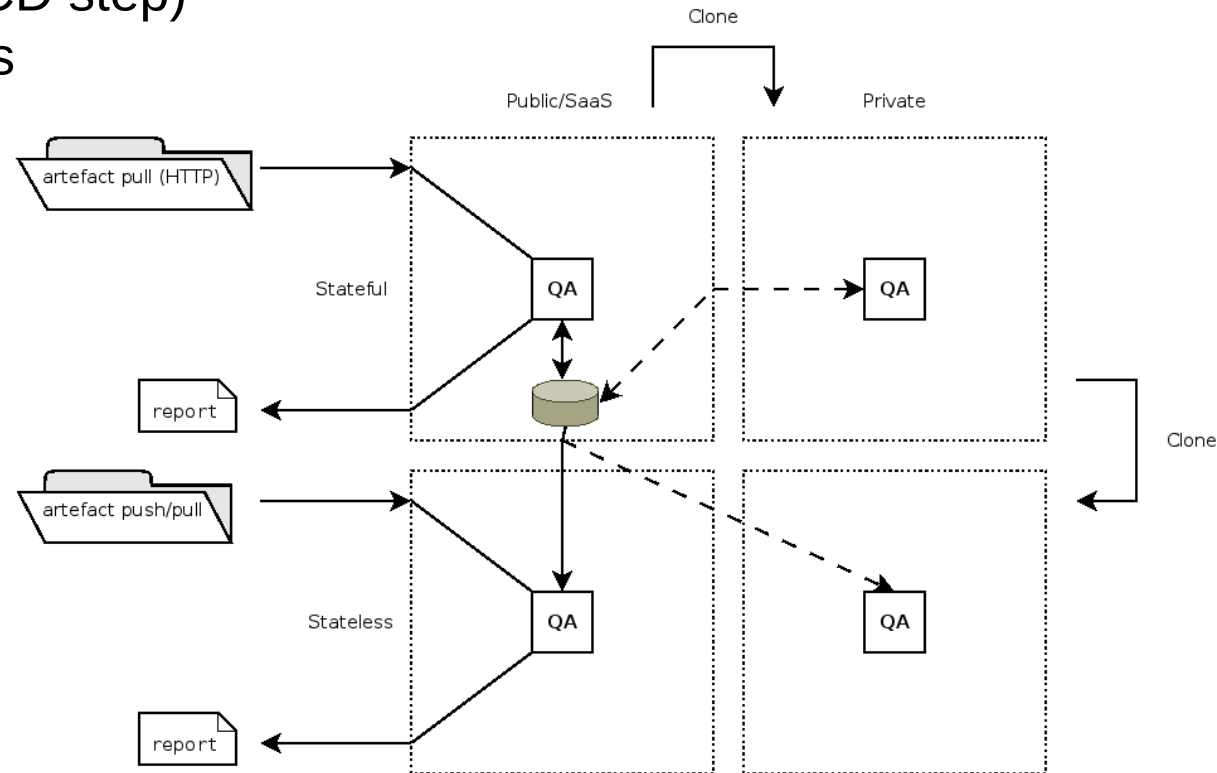
Researcher interest

- regular batch generation of results
- not time-critical
- holistic view including history

Hybrid assessment

- important for regressions

→





➤ ➤ Rewind



# Key points

## Summary

- Software development mindset: code/artefact/(micro)service levels
- Quality assessment perspective
- Temporal perspective - software (artefact) evolution + quality evolution

## Dual purpose effort

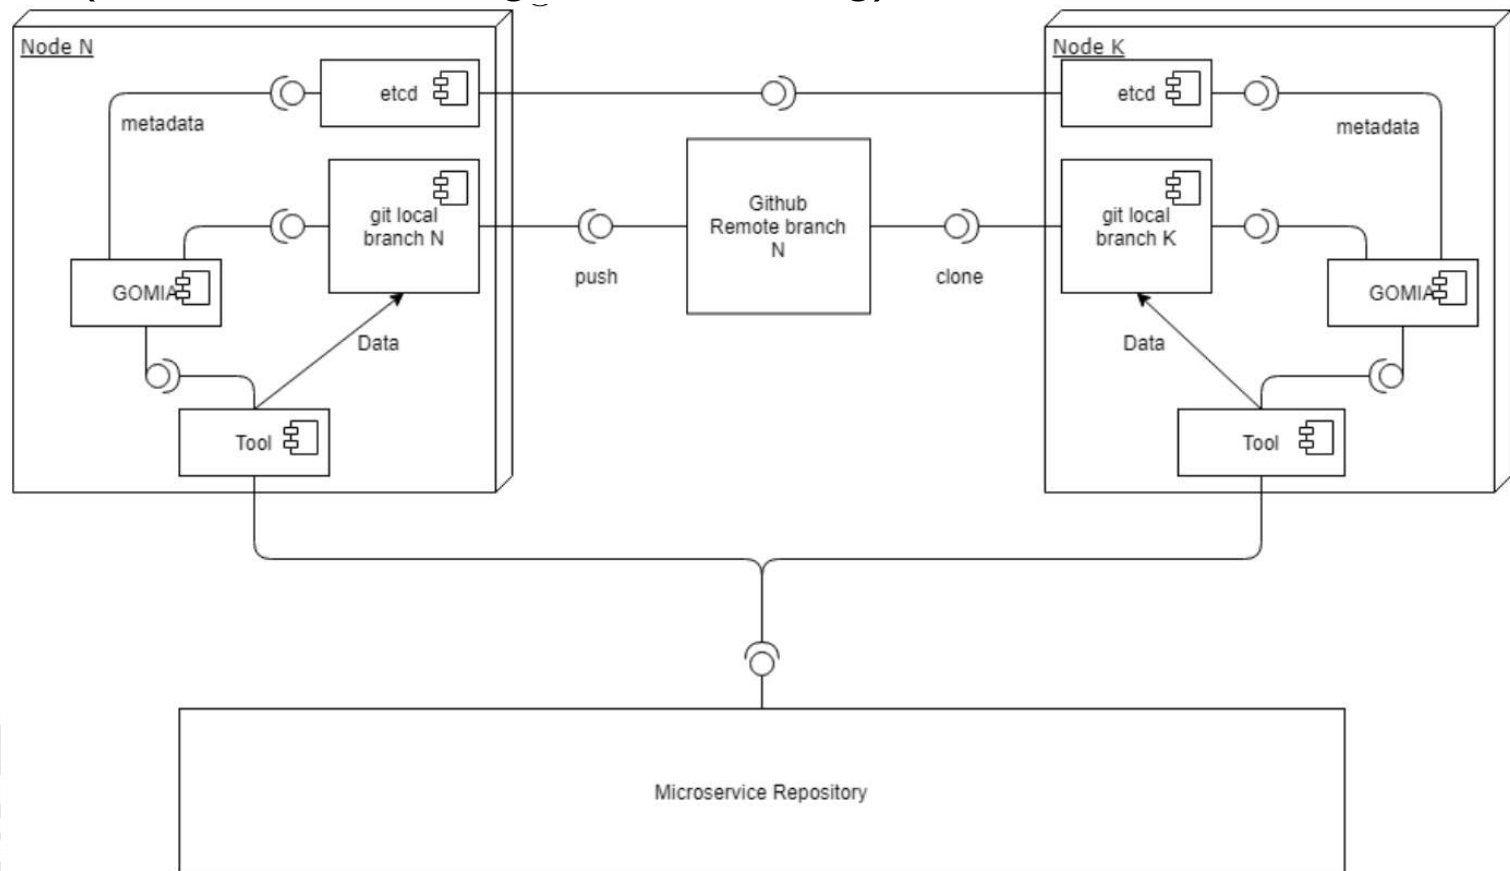
- practical tools for developers
- insightful research infrastructure for scientists



# Research on artefact quality+evolution

MAO = Microservice Artefact Observatory

Artefact type-specific checkers (tools) + decentralised architecture + smartness (resilience, self-regression testing)



# Research results

## HelmQA checker

HelmQA by Service Prototyping Lab, Zurich University of Applied Sciences  
Follow the links to per-chart quality reports.

[aerospike-0.1.7.tgz](#)  
[anchore-engine-0.1.7.tgz](#)  
[artifactory-5.4.6-3.tgz](#)  
[artifactory-7.1.5.tgz](#)  
[artifactory-ha-0.1.9.tgz](#)  
[bitcoind-0.1.3.tgz](#)  
[buildkite-0.2.3.tgz](#)  
[burrow-0.4.3.tgz](#)  
[centrifugo-2.0.1.tgz](#)  
[cert-manager-0.2.9.tgz](#)

HelmQA advice on chart traefik-1.1.0-a.tgz / data point 2018-05-31

**Issue: multiple:traefik-1.9.0.tgz,traefik-1.1.2-h.tgz,traefik-1.2.1-b.tgz,traefik-1.1.1-a.tgz**

Explanation:

Two or more versions of the same chart appear in the same repository.

Rationale:

Developers may not understand which version to use.

Recommended action:

Unless there are compelling technical reasons, the advice is to stick to a single version per category.

**Issue: duplicatevalues**

Explanation:

The templates contain values with at least two duplicates which are candidates for template variables.

Rationale:

Configuration duplication should be avoided. While false positives may occur, in some cases the duplication is intentional.

Recommended action:

Review the list (and if available, the suggested diff) to decide which duplicates should be removed. The following occurrences of duplicate values have been found. Each entry consists of the value and the number of occurrences.  
[[ '80', 6], [ '10', 4], [ '1', 4], [ 'TEMPLATE-dashboard', 3]]

[diff sketch available](#)



# Research results

Docker Compose checker

Docker Compose Validator!



Github Repository



Upload Docker-  
compose

✕ Filters



Duplicate service name



Duplicate container name



Duplicate image

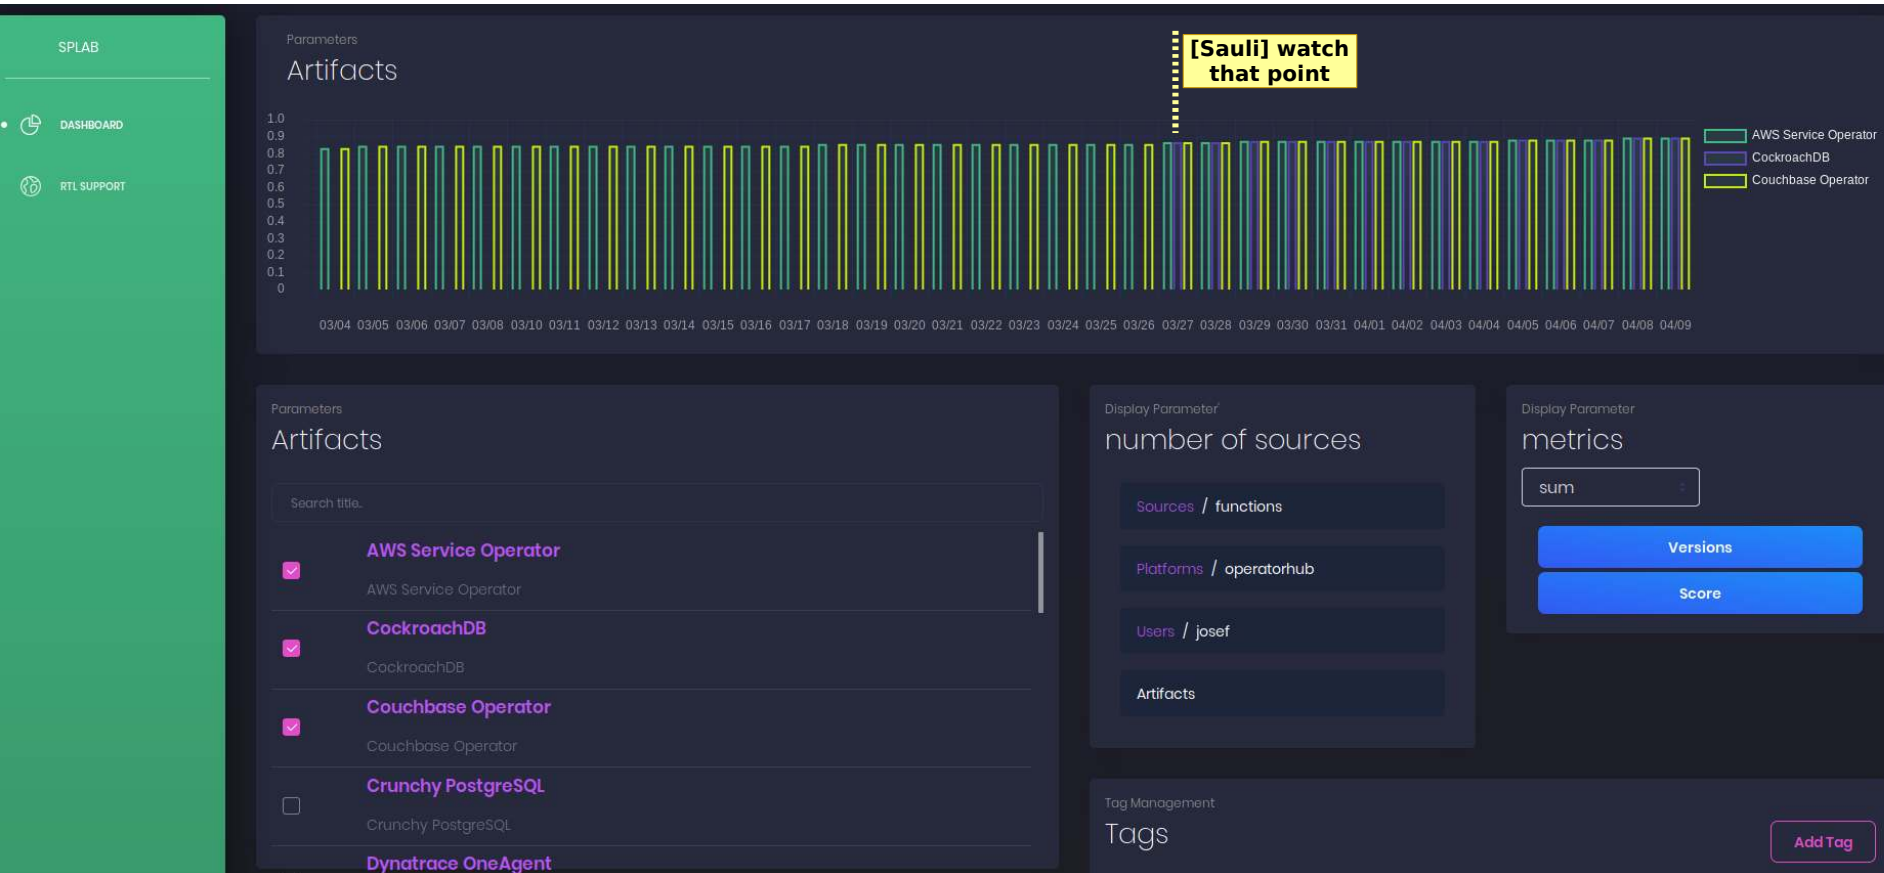


Duplicate port



# Research results

## Observatory dashboard across artefact types



- specialised roles (developer, observatory operator, statistician)
- flagging & debating research data “points of interest” with peers/citizens

