

Cloud service dependencies in practice


Josef Spillner <josef.spillner@zhaw.ch>
Service Prototyping Lab (blog.zhaw.ch/splab)

Sep 04, 2019 | INFORTE Summer School, Tampere

I. Dependencies Primer



“Standing on the shoulders of giants”




Welcome to reuse heaven



Welcome to dependency hell

Overview about dependencies

Abstract considerations



Distinguish:

- static vs. dynamic
- visible/recognisable vs. hidden
- explicit vs. implicit
- strict vs. tolerant / strong vs. weak



System model: deps & artefacts

Dependencies in software

- digital artefacts
 - code (executable) - also: components, services
 - configuration (declarative)
 - data (binary)
- specialisations of software engineering process
 - CBSE - component-based
 - SOSE - service-oriented
 - microservice-oriented...



Compositional model

“Compositio” and “Componentens”

Composition

- of multiple components
- some equal, some with distinct function
 - homogeneous
 - heterogeneous
- goal: whole $> \Sigma$ parts



[citysuburb.wix.com]

Component

- piece of hardware (out of scope)
- piece of software
- restricted view
 - black box: only interface
 - grey box: partial information about the inside and behaviour



Roles: Each object can be component and composition at the same time.



Composition of services

Scenario: New mobile app for mensa food delivery to lecture room



- registration e-mail: re-use existing software?
- menu selection: re-use current menu listing?
- payment: handle different methods?

→ main drivers: re-use, extensibility



Service dependencies

Importance

- if B goes down, A goes down, too
- hidden cost

Dependency manifestation

- tight coupling / fixed binding
- loose coupling / late binding

Dependency type

- inter-service
- intra-service (e.g. libraries, backend services)

Dependency consideration

- data sources
- transitivity

A: application (service)

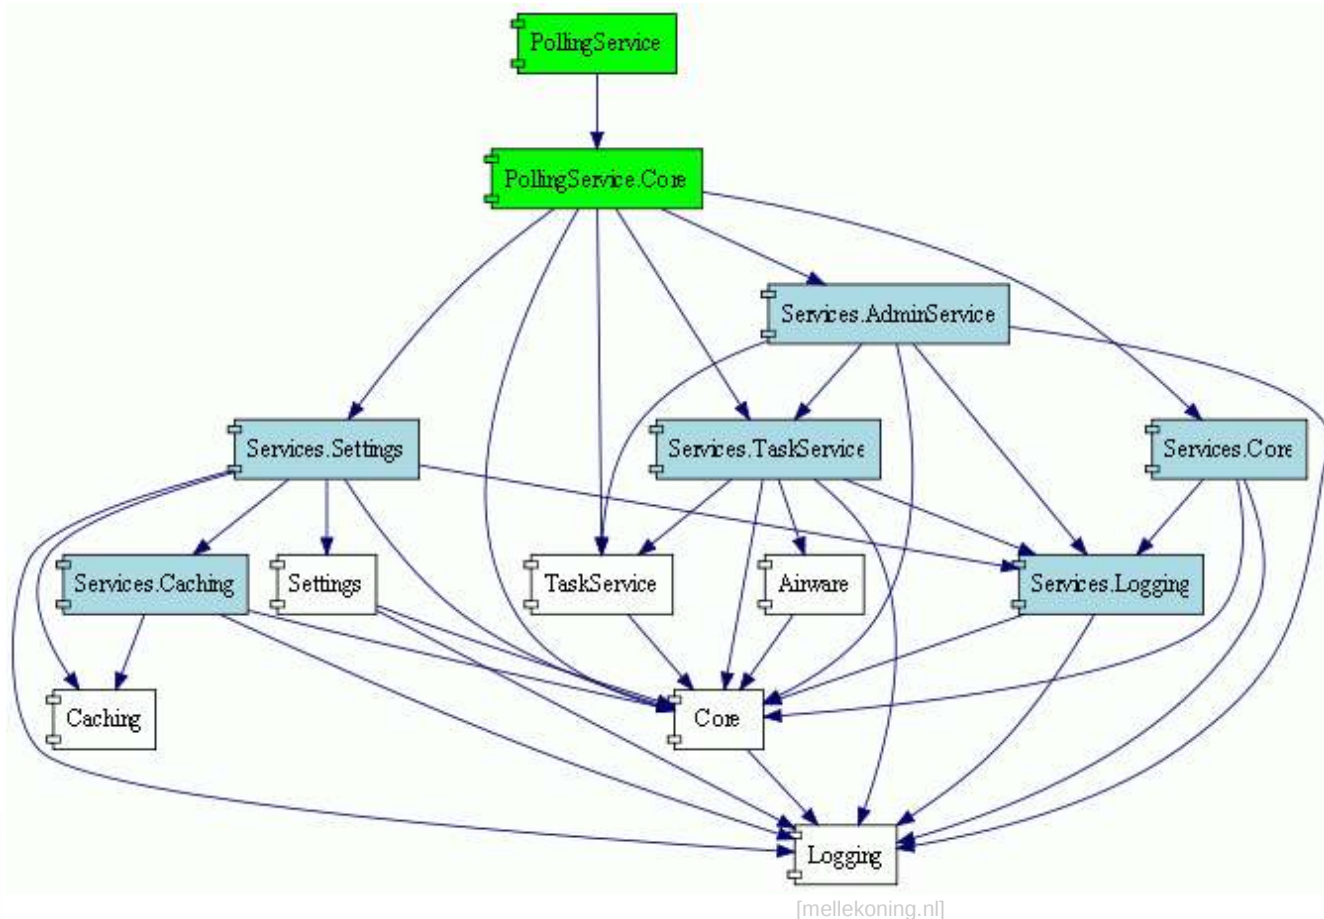


B: external service



Dependency graphs

directed, cyclic (to be avoided) or acyclic, weighted or unweighted graph



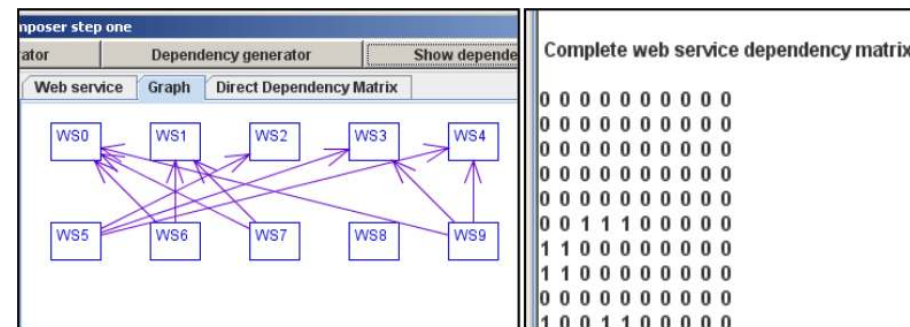
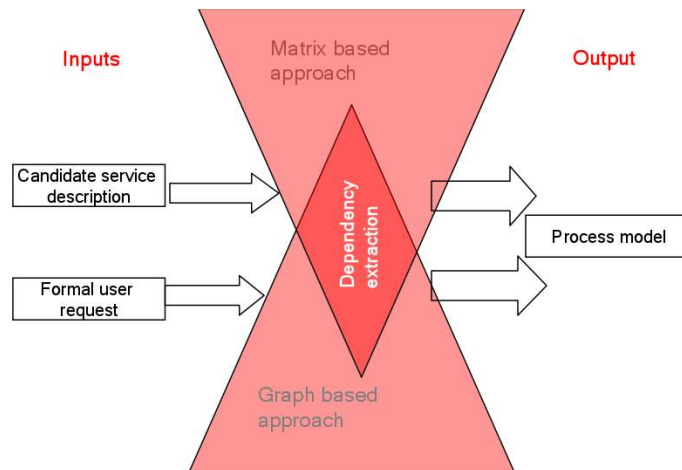
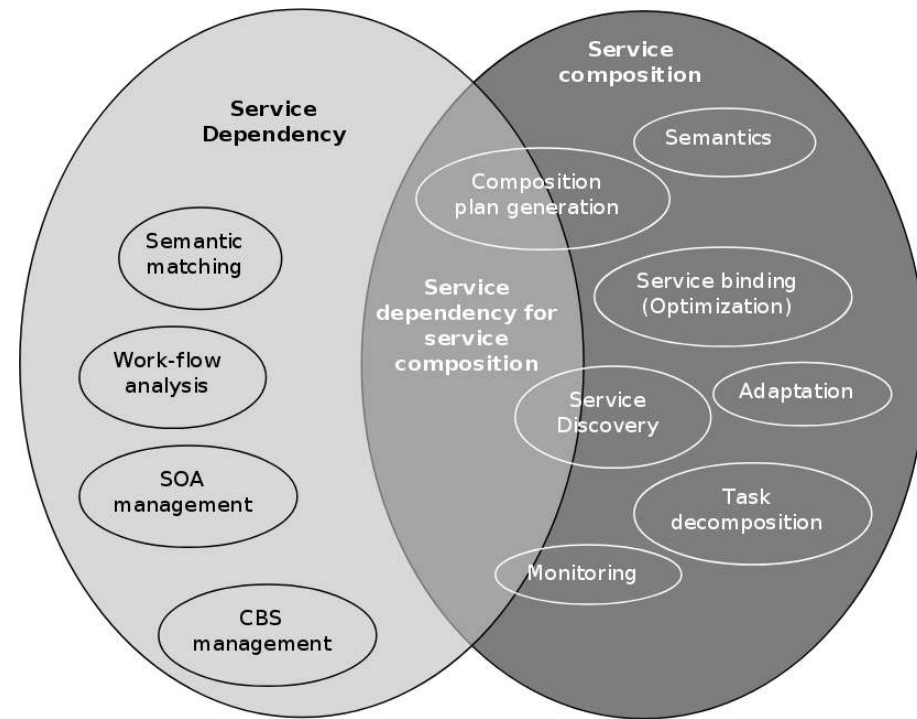
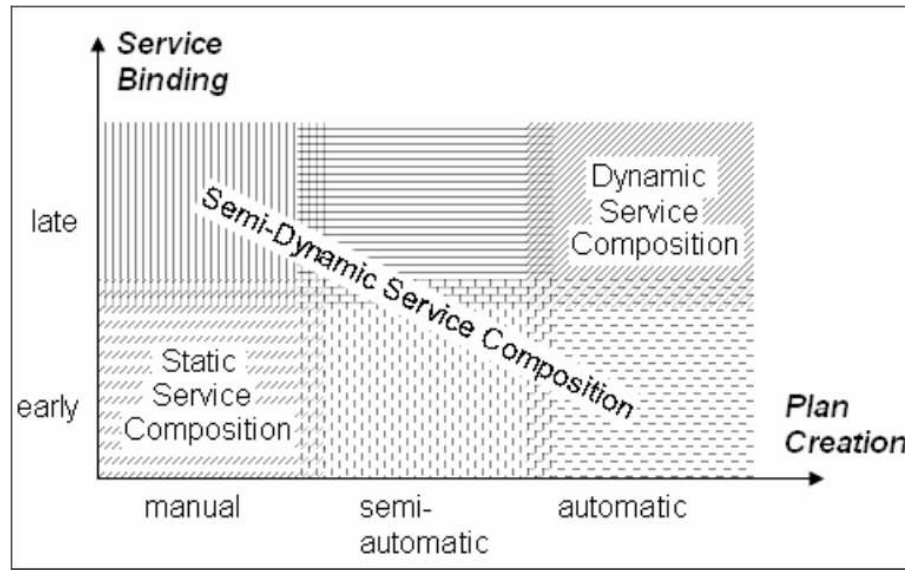
[mellekoning.nl]



Dependency graphs

Service composition classification + automated creation

[Omer'11]




II. Detection/Resolution



Dependency analysis

A depends on B if

- declared as such
- A invokes B (traceable)
- logic of A depends on state within B (not traceable)



Analysis techniques

- statically on description
- dynamically during execution

Dynamic analysis requirements

- correlation vs. causation (coincidental vs. true)
- small data sets (infrequent/conditional use)
- clustering of services (e.g. backup/failover)



Dependency analysis after Peddycord

[USENIX LISA'12]

3 novel identification techniques

basis:

- passive network monitoring and analysis
- logarithm-based ranking scheme
- frequency inference


implementation: extended NSDMiner (NSD: Network Service Dependency)

model: communications graph

heuristic calculation: confidence in dependency candidates

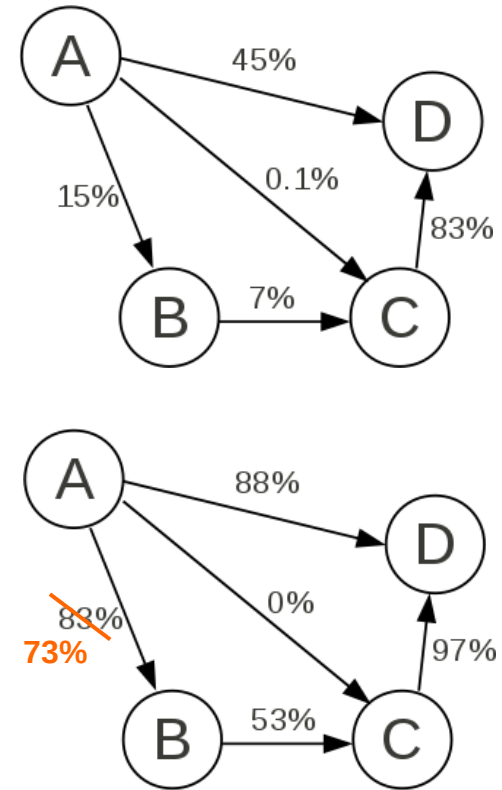


Dependency analysis: example



ratio

logarithm



node weights: #occurrences as destination
edge weights: #nested occurrences

confidence as ratio-based ranking: $\text{weight}(A \rightarrow B) / \text{weight}(A)$

confidence as logarithm-based ranking: $\log_{\text{weight}(A)} \text{weight}(A \rightarrow B)$



Dependency resolution

Immediate resolution

- simple transitive inclusion
- solution not guaranteed

Interactive resolution

- with feedback
- iterative improvements
- may complement immediate resolution upon escalation

SAT (satisfiable) solver resolution

- a solution which exists will be found
- generation of proof for unsolvable problems
- NP-complete calculation

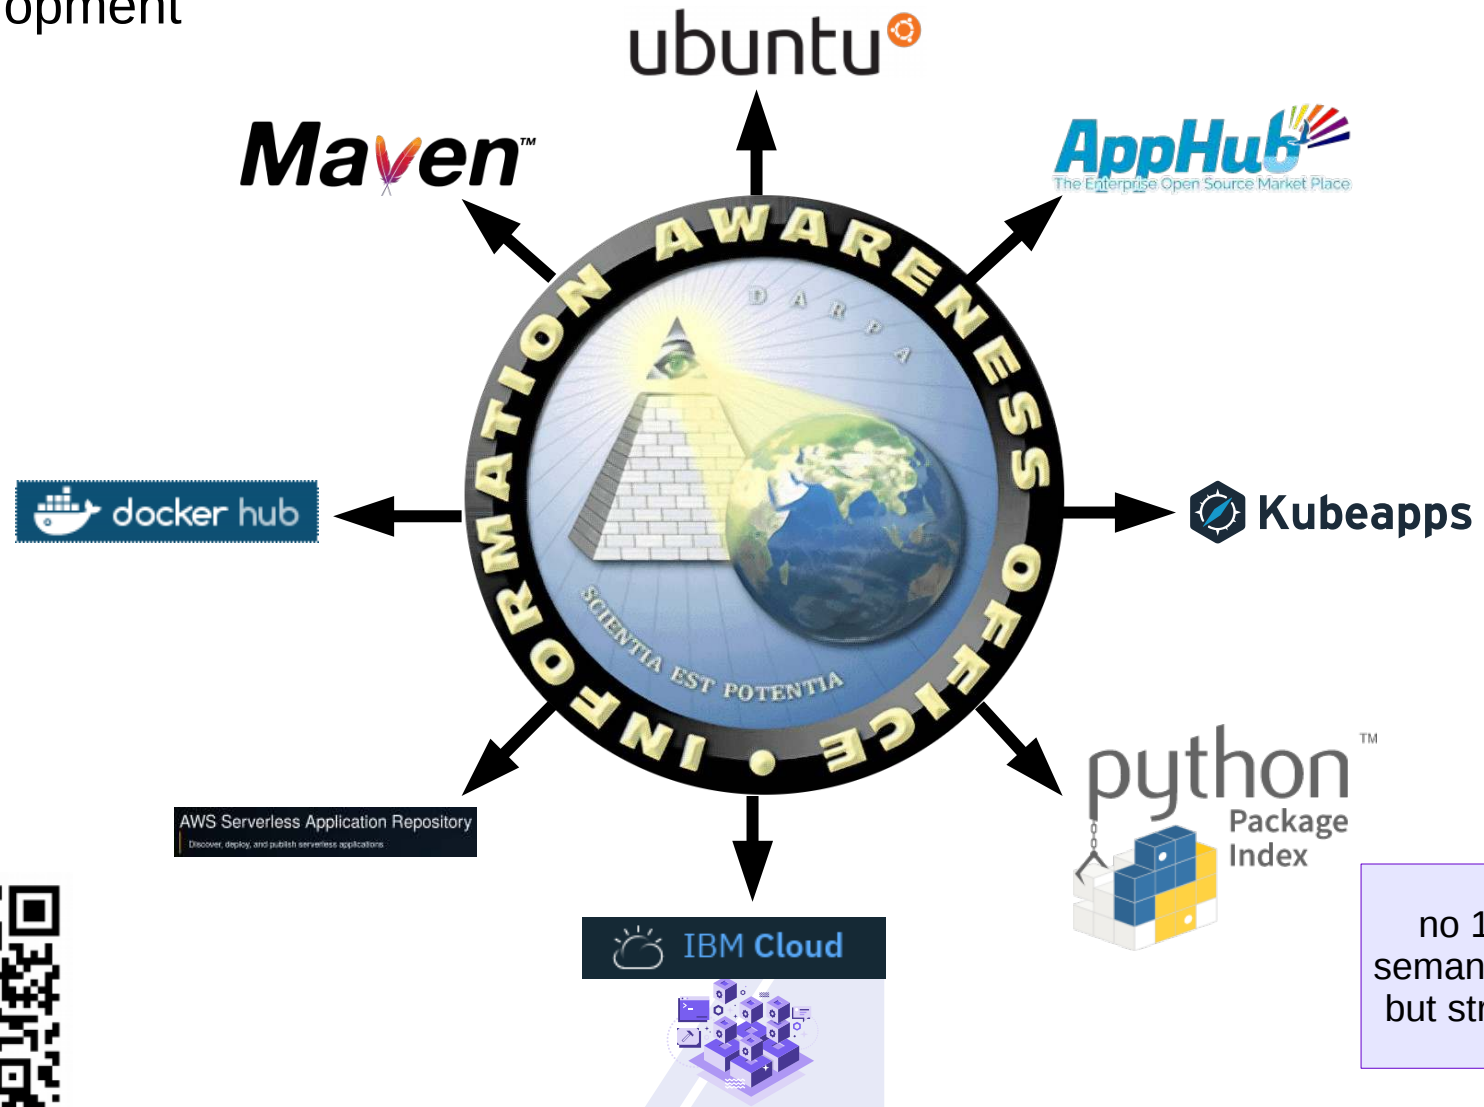


III. Composite Applications



Practical service dependencies

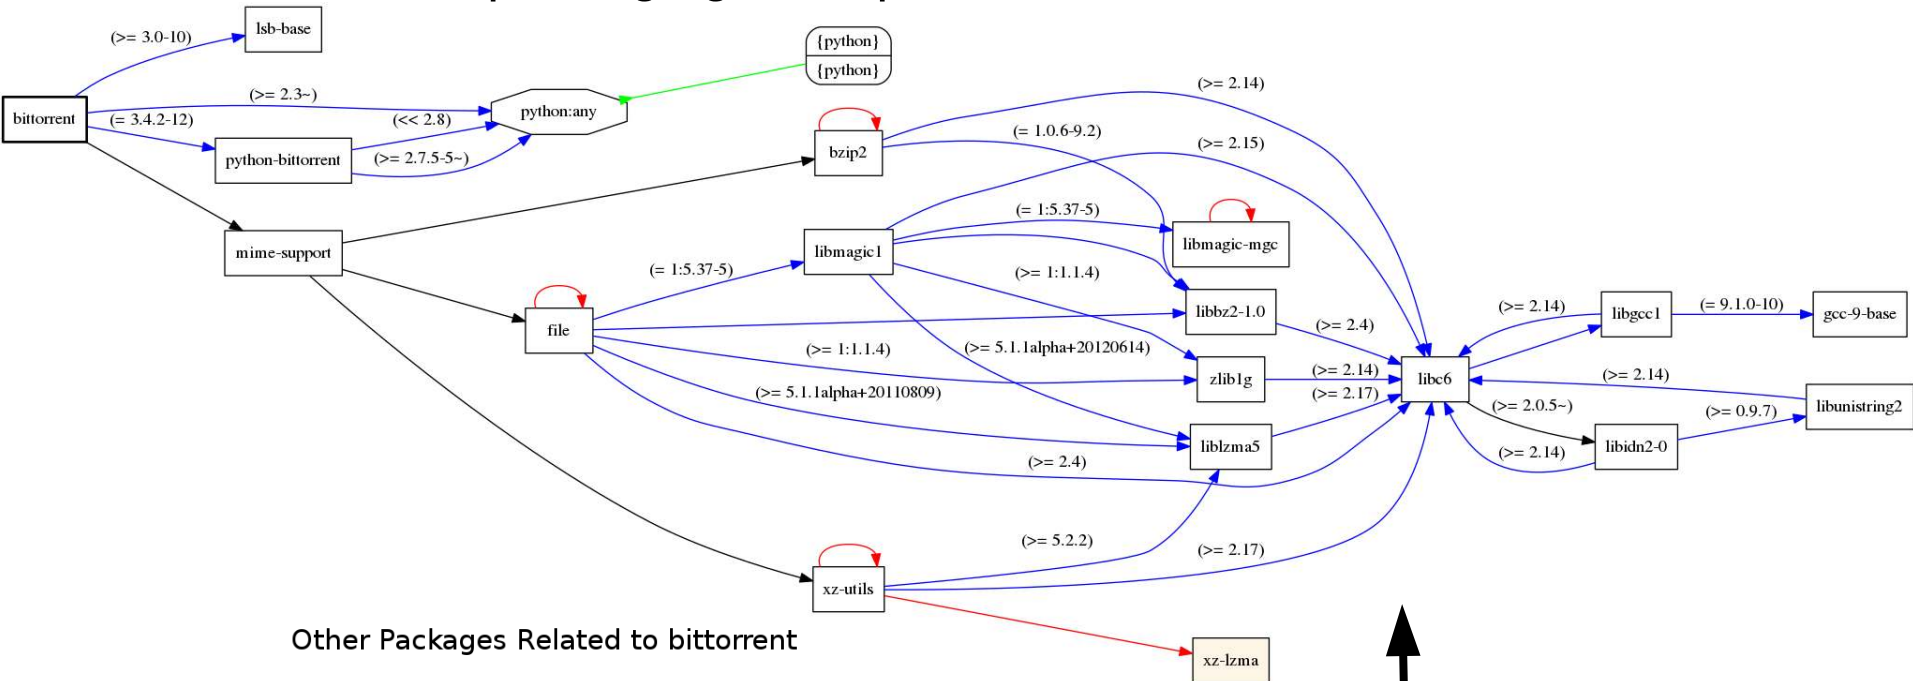
Artefact repositories: The backbone of today's (cloud) software development



no 1:1 service semantics mapping but strong relation

Basis: software dependencies

Linux distribution packaging example: Bittorrent



Other Packages Related to bittorrent

• depends ■ recommends ◆ suggests ▲ enhances

- **dep: [lsb-base](#) (>= 3.0-10)**
Linux Standard Base init script functionality
- **dep: [python](#)**
interactive high-level object-oriented language (Python2 version)
- **dep: [python-bitTorrent](#) (= 3.4.2-12)**
Scatter-gather network file transfer
- **rec: [mime-support](#)**
MIME files 'mime.types' & 'mailcap', and support programs
- ◆ **sug: [bittorrent-gui](#)**
Original BitTorrent client - GUI tools

○ Build-Depends: dark gold, bold
 ○ Build-Depends-Indep: light gold
 ○ Pre-Depends: purple, bold
 ○ Depends: blue
 ○ Recommends: black
 ○ Suggests: black, dotted
 ○ Conflicts: red
 ○ Provides: green, inverted arrow

○ not applicable to service dependencies

Practical service dependencies

Doodle example

Monatliche Sitzung
 Umfrage von Hans | 1 | 0 | vor weniger als einer Minute

Wann haben Sie Zeit für unsere monatliche Sitzung?

Tabellen-Ansicht | Kalender-Ansicht

Dies ist eine Beispiel-Terminumfrage.

Mehr erfahren ...




1 Teilnehmer

	November 2015					
	Mo 9	Di 10	Do 12			
12:00 – 14:00						
09:15 – 11:15		✓		✓	✓	
14:45 – 16:45						
	0	1	0	1	1	

Ich kann nicht | Speichern

HTTP GET/POST
improper REST

JSF scaffold
JSON data



SMTP



POSTFIX



JSON documents

HTTP
ICS/
iCal



Kolab.org



3rd party
services



Practical service dependencies

Restlet-based applications: dependencies on .ext.atom, .ext.gae, .ext.sip etc.

- expressed in pom.xml files

Note: No generic standard notation available

Circular/recursive dependencies: should be avoided, but...

PAGE 3

DEPARTMENT	COURSE	DESCRIPTION	PREREQS
COMPUTER SCIENCE	CPSC 432	INTERMEDIATE COMPILER DESIGN, WITH A FOCUS ON DEPENDENCY RESOLUTION.	CPSC 432



Composite services (code deps)

Characteristics

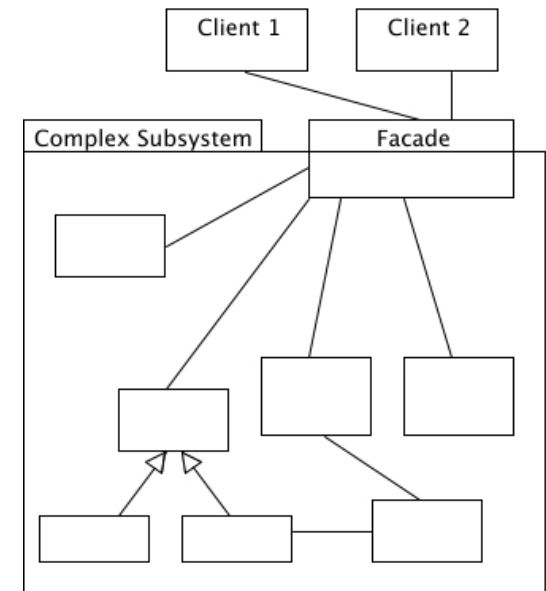
- offer a single service interface
- distribute requests to multiple services within the composition (in parallel, serially, or more complex routing)
- require knowledge of dependencies (internally or by the caller)

Remember software engineering – design patterns

- Composite / Façade
- Factory method

Advantages for services

- improved QoS – e.g. higher availability
- improved QoE – e.g. flexibility to switch



Composition techniques/patterns


- Orchestration
- Choreography
- Bundling
- Multiplexing
- Mashup (out of scope)
- Service Bus (out of scope)



Orchestration


Service Orchestration:

- multiple services form another one
- centrally managed and executed
- based on workflows
- workflows require service interfaces



Resource Orchestration (e.g. for cloud computing):

- workflows require resources (e.g. software implementation of service)
- resource allocation must be performed centrally



Service orchestration

Definition: coordinated arrangement of service invocations; may be executable as another service

Potential benefits: creation of value-added services by re-using others

Potential risks: issues with dependency services (unavailable, faulty)

Example excerpt (in DSOL):

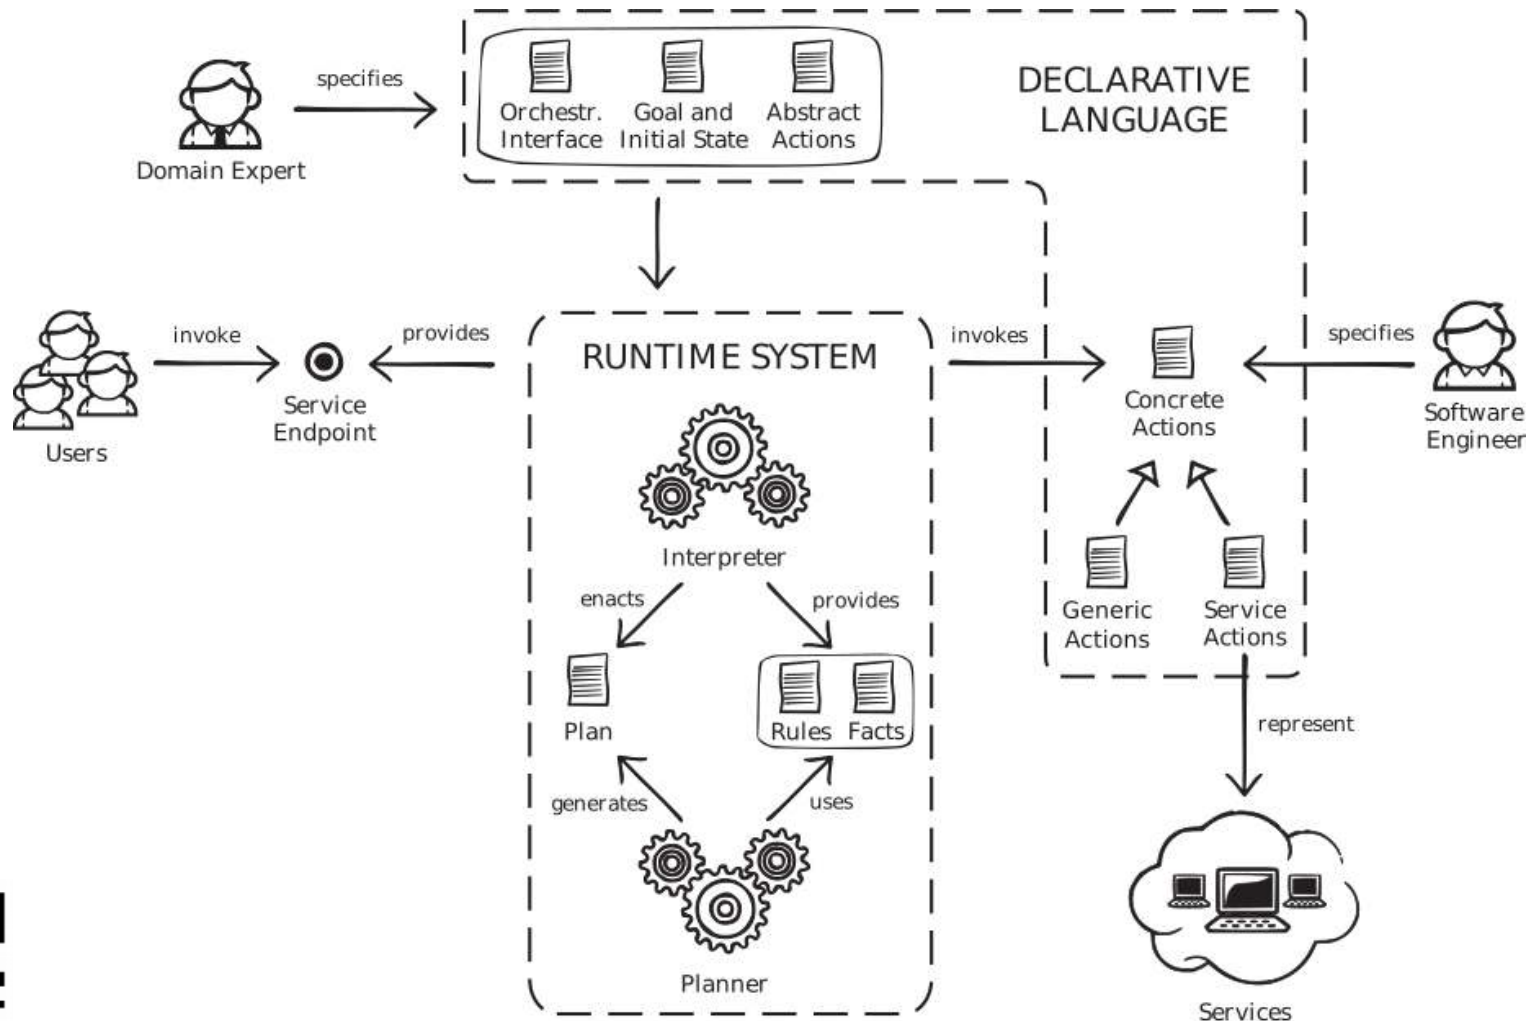
```
<services>
<service type='rest' name='poll' id='doodle' method='GET' url='.../{pollId}' />
<service type='soap' name='translate' id='translate' operation='Detect' wsdl='...' />
</services>
```

```
@WebService
public interface PollTranslator {
    @WebResult(name='translatedPoll')
    public Poll getTranslatedPoll(@WebParam(name='pollId') String id,
                                @WebParam(name='language') String desiredLanguage);
}

...
```



Orchestration with DSOL - example



Orchestration languages

Service Orchestration

- Business Process Execution Language (BPEL)
- Yet Another Workflow Language (YAWL)
- Dynamic Service Orchestration Language (DSOL)
- Workflows
 - Montage etc.
- Cloud function workflows
 - AWS Step Functions
 - Fission Workflows
 - IBM Composer

Resource/Implementation Orchestration

- Heat Orchestration Template (HOT)
- AWS CloudFormation
- Vamp Blueprints + Workflows
- Docker Compose
- Kubernetes Descriptors
- Juju Charms

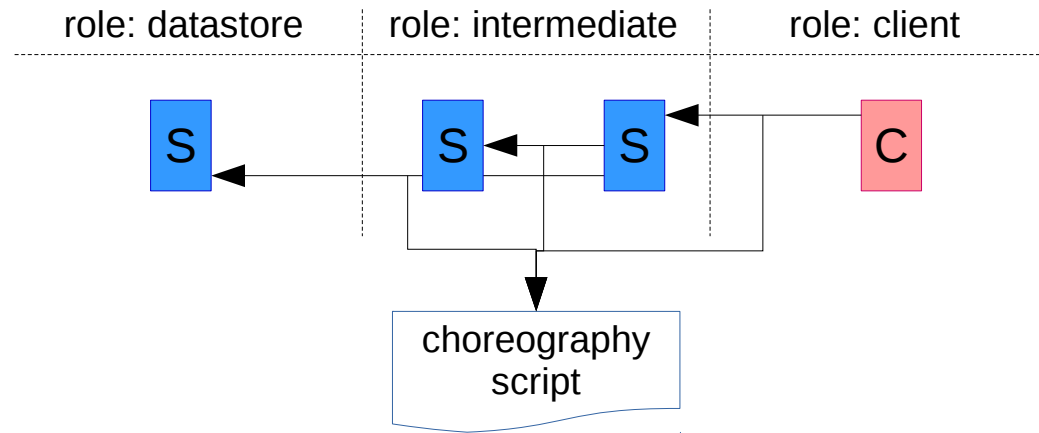


Choreography

Definition: global interaction protocol between autonomous service partners

Potential benefits: no central point of control; declarative messaging behaviour

Potential risks: difficult decentralised enactment; little industry acceptance



Languages: WS-CDL (historic), BPMN 2.0, Chor, CHOReVOLUTION




Bundling

Definition: multiple services offered/used in a bundle

(Mixed bundles: service access + tools, clients, other products)

Potential benefits: cheaper, less administrative overhead

Potential risks: less flexibility for exchanging single service



Languages: USDL




Multiplexing

Definition: multiple services used in parallel handling partial requests

(compare load balancing or failover: serial use)

Potential benefits: flexible redundancy schemes, “survival of the fittest”

Potential risks: more administrative overhead, higher cost due to candidates



Languages: none



Composite microservices

Definition: single-function-oriented services which scale elastically and operate resiliently backed by a portable implementation.

Inherited definition from arbitrary services:

- loosely coupled
- uniformly described and invoked
- composable
- re-usable

Implementations: containers, hosted functions, unikernels...



Microservices management platforms

Tasks

- deployment and management of microservice representations (i.e. containers)
- partial upgrades without downtimes, honouring dependencies
- canary testing
- monitoring, migration, scaling, ... without downtimes
- ... and: rapid prototyping!

Implementations (selection)

- Vamp
- Kubernetes / OpenShift
- Apache OpenWhisk
- provider-specific: AWS ECS, Azure CS, Google CE, IBM CS for Bluemix, ...



APACHE
OpenWhisk™



vamp

Composition & management example

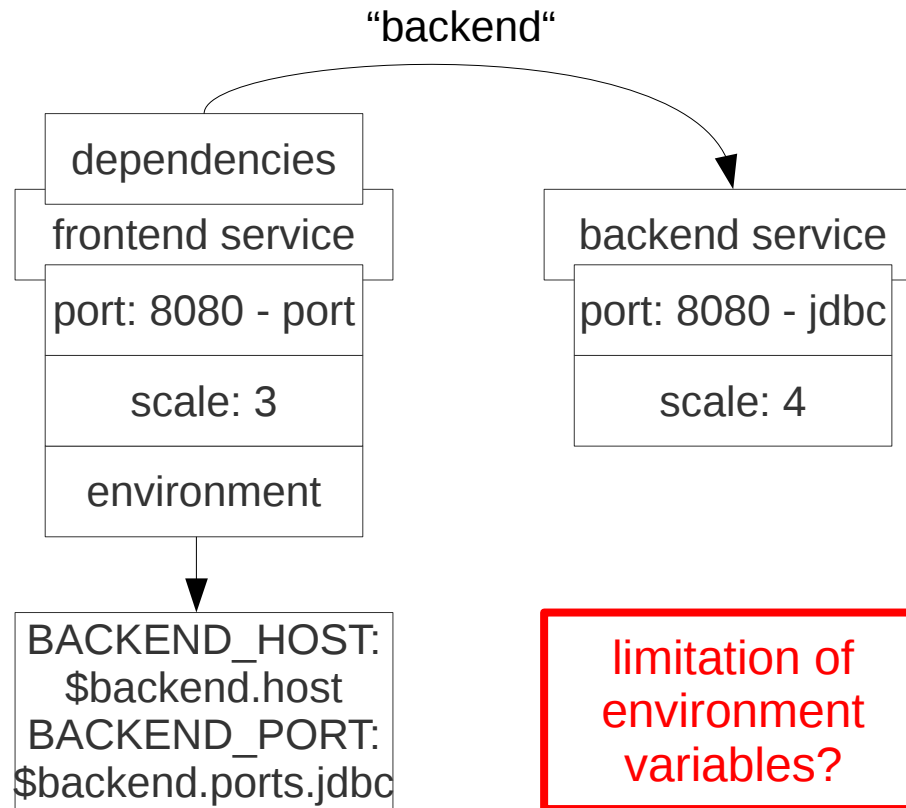
Vamp

- using “blueprints”
- ingredients: services, clusters, gateways, conditions, escalations, etc.

```
---
name: my_blueprint                # Custom blueprint name
gateways:
  8080/http: my_frontend/port
clusters:
  my_frontend:                    # Custom cluster name.
    gateways:                     # Gateway for this cluster services.
      routes:                     # Makes sense only with
        some_cool_breed:          # multiple services per cluster.
          weight: 95%
          condition: User-Agent = Chrome
        some_other_breed:         # Second service.
          weight: 5%
    services:                     # List of services
      -
        breed:
          ref: some_cool_breed
        scale:                     # Scale for this service.
          cpu: 2                   # Number of CPUs per instance.
          memory: 2048MB           # Memory per instance (MB/GB units).
          instances: 2             # Number of instances
```



Service discovery example

Again, in Vamp:



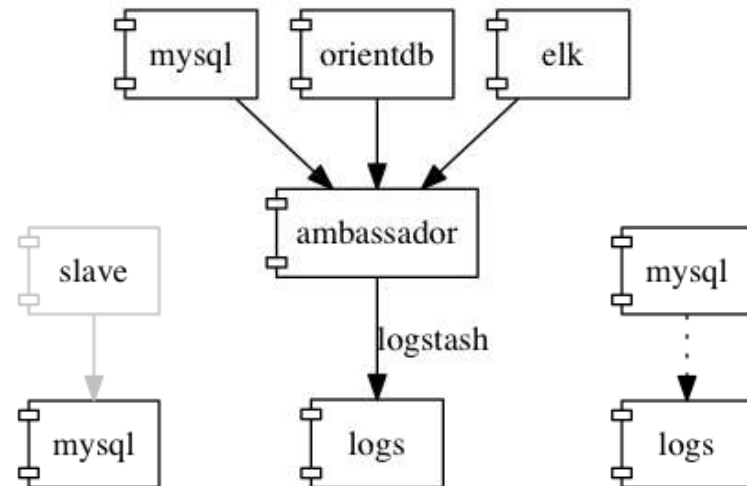
Dependencies in Docker Compose

Dependencies are *by reference* → need resolution before execution



Internal + external service dependencies: links + deps

Volume dependencies



Issues with deps in Docker Compose

- Docker images
 - do not (any longer) exist
 - are not accessible (i.e. registry requires login)
 - brings along defect or security vulnerability
- Volumes
 - do not not (any longer) exist
 - are not accessible (permission issues)
 - are full
- External services
 - are not accessible (networking level, authorisation)
 - are not yet ready → wait scripts like wait-for-it.sh



Dependencies in Helm charts

Dependencies are *by embedding* - no resolution needed - but: *transitivity*!

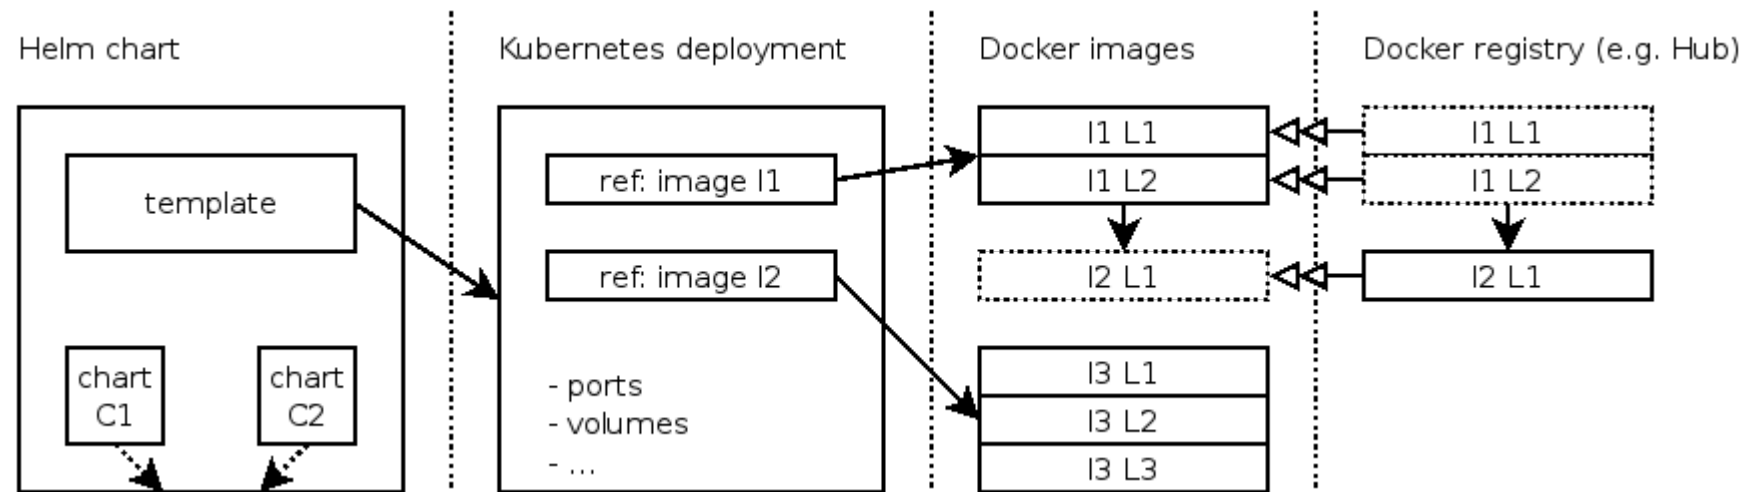
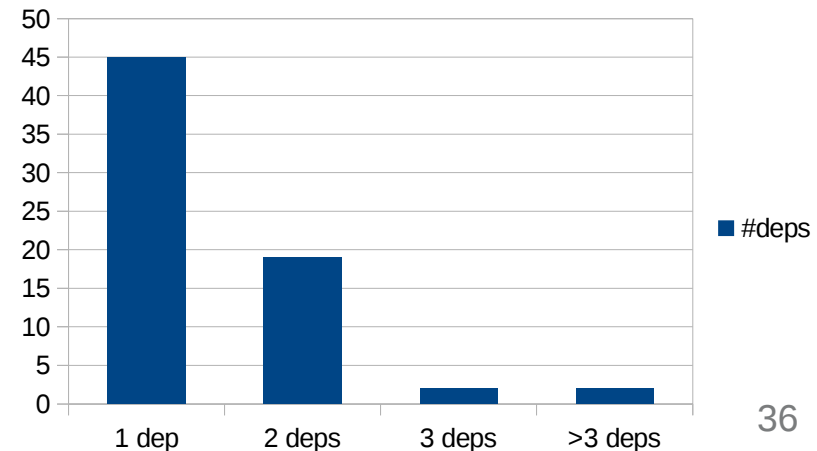


Chart structure →

```
chartname/Chart.yaml
chartname/charts/dep1/Chart.yaml
chartname/charts/dep2/Chart.yaml
```

23% of all charts (on KubeApps Hub) have dependencies; distribution [Jul'19] →



Issues with deps in Helm charts

Issues:

- in 0.3% of all cases: incorrect dependencies (chartname/dep1/Chart.yml in cert-manager)
- among all charts with deps:
 - 23.5% have incompletely specified dependency versions (e.g. sugarcrm, quassel, sonarqube); among the remainder:
 - some are fine (e.g. elasticsearch-exporter: provides 1.0.2; elastic-stack requires 1.0.2)
 - some are not so fine (e.g. mongodb: provides 4.0.11; charts like rocketchat require 4.0.10, 4.0.9, 3.6.4) - totalling 37.2%
 - subject to bitrot, vulnerabilities, deprecations, failure to resolve docker images...

Issue score: $0.3\% + 23\% \cdot 23.5\% + 76.5\% \cdot 37.2\% = \mathbf{34.2\%}$ of all Helm charts can't do deps



Application Portability




Portability and migration

Portability $\leftarrow \rightarrow$ dependency of application/service on platform

Migration $\leftarrow \rightarrow$ resolving dependency

Complexity of migration

- homogeneous
- inhomogeneous
- heterogeneous



Types of microservice migrations

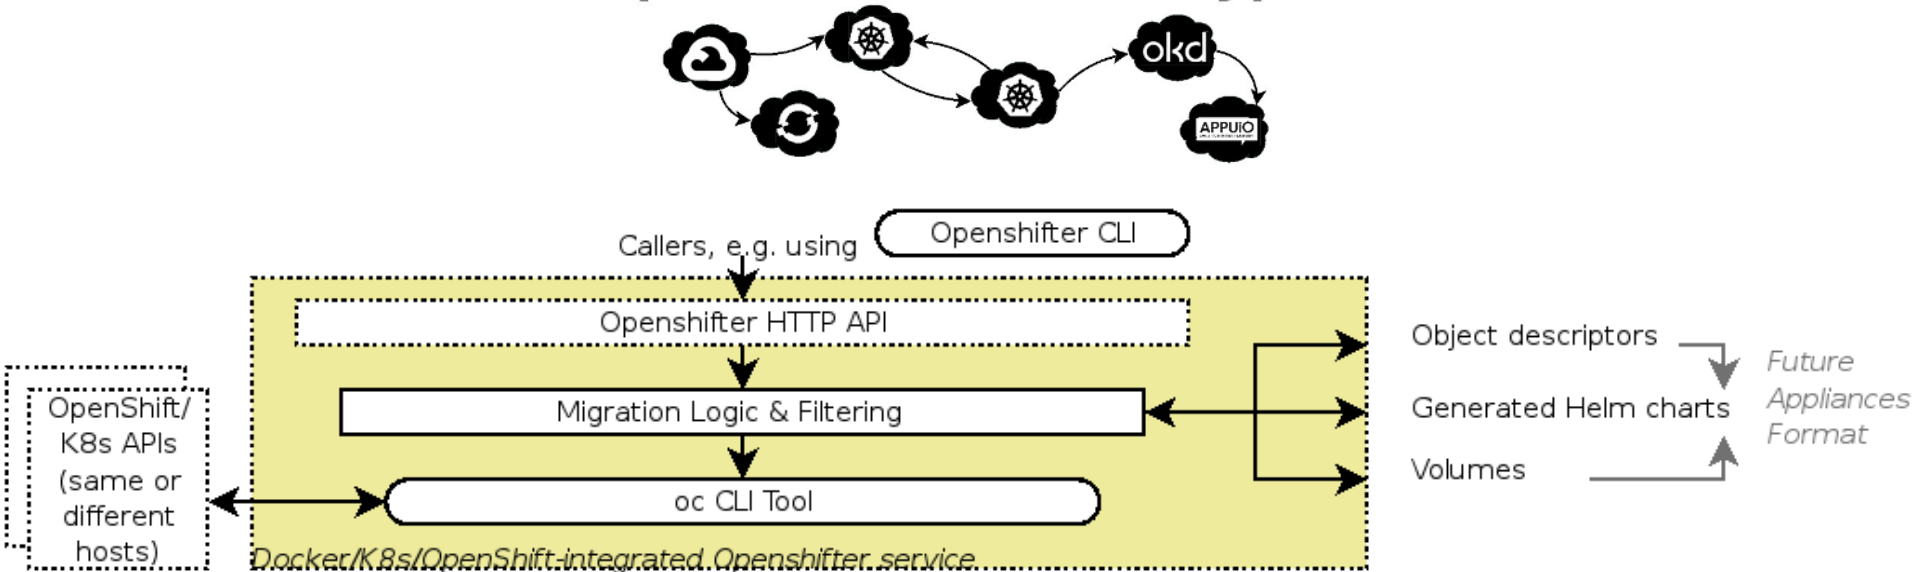
- onboarding: dev \rightarrow prod
- region change: prod1 \rightarrow prod1
- upgrades (e.g. storage): prod1 \rightarrow prod1'
- external backup: prod \rightarrow backup



Scenario: OpenShifter

Migration of OpenShift application + volumes

Openshifter Prototype



```
curl https://localhost:8443/export/<baseURL>/<project>/<user>/<password> >
_output
base64 -d < _output> _output.tgz
```

```
curl https://localhost:8443/delete/<baseURL>/<project>/<user>/<password>
```

```
curl -X POST --data-urlencode @_output.tgz https://localhost:8443/import
/<baseURL>/<project>/<user>/<password>
```


Scenario: FaaS Converter

Migration of cloud functions

- entry point function signature differences
 - faasconverter
- management API (control plane) differences
 - snafu-import

AWS Lambda:	OpenWhisk/IBM Functions:	OVH Functions:
<pre>def lambda_handler(event, context): """ event: dict context: meta information obj returns: dict, string, int, ... """ # ... return "result"</pre>	<pre>def handler(input): """ input: dict returns: dict """ # ... return {}</pre>	<pre>def handler(input): """ input: dict returns: str """ # ... return ""</pre>
Fission:	Azure Functions:	Further differences:
<pre>def main(): """ input: flask.request.get_data() returns: str """ # ... return "result"</pre>	<pre>def main(): from AzureHTTPHelper \ import HTTPHelper input = HTTPHelper().post # ... open(os.environ["res"], "w").\ write(json.dumps({"body": "..."})) main()</pre>	<ul style="list-style-type: none">• function naming (mangling on client or service side)• function granularity (number of entry points)

```
./faasconverter --file test.py --function foo --providers aws, azure --just-wrap False --all-together True
```

Options

file File to convert to the selected provider syntax

function Selected function to convert to the providers syntax

just-wrap To only add the syntax wrapper on the end of the selected file

providers List of selected providers to which to convert the functions, available = aws, ibm, azure, ovh, fission

all-together Put all the wrappers together and add it to the the end of each file



➤ ➤ Rewind



Summary

Dependency lifecycle & representation

- extraction/generation/manual specification + analysis/identification + resolution + optimisation/adaptation
- graph, matrix

Complexity of dependency considerations

- services vs. software artefacts in service-oriented compositions
- variety: static vs. dynamic, visible vs. hidden, strong vs. weak

Duality with compositional models

- orchestration, choreography, bundling, multiplexing
- applicability to composite microservices (& artefacts)

